

Multi-Agent Tool-Integrated Policy Optimization with Process Reward

Zhanfeng Mo^{* 1} Xiaogang Xu^{* 1} Xingxuan Li^{* 1} Lidong Bing¹

Abstract

Large language models (LLMs) increasingly rely on multi-turn tool-integrated planning (TIP) to solve open-ended, knowledge-intensive, and complex reasoning tasks. Existing single-agent TIP frameworks suffer from limited context length and noisy tool responses, motivating the use of multi-agent frameworks with planner- and worker-agents to manage context. However, effective reinforcement learning (RL) post-training methods for tool-integrated multi-agent frameworks remain underexplored. To address this gap, we propose Multi-Agent Tool-Integrated Policy Optimization (**MATPO**), a principled policy gradient method that enables distinct roles (planner and worker) to be trained within a single LLM instance using role-specific prompts. Furthermore, we develop **MATPO-PR** (MATPO with Process Reward) to tackle the credit assignment problem in multi-turn planner-worker interactions by optimizing the quality of each intermediate worker-agent response. Our experiments show that MATPO-PR consistently outperforms single-agent baselines in both performance and robustness to noisy tool responses on three deep research benchmarks. Our findings highlight the effectiveness of unifying multiple agent roles within a single LLM and provide practical insights for effective RL training with multi-agent TIP systems. Our code is available at <https://github.com/xiaogang00/MATPO-PR>

1. Introduction

Tool-augmented LLM agents have revolutionized workflows across domains by enabling LLMs to solve complex problems through external tool use (Dong et al., 2025; Qian et al., 2025; Qu et al., 2025; Li, 2024; Singh et al., 2025; Chen et al., 2025a). These agents operate via multi-turn tool-integrated planning (TIP): after receiving a user query, an agent performs multiple iterations of planning-and-tool-calling until it is confident enough to produce a final answer.

^{*}Equal contribution ¹MiroMind AI, Singapore. Correspondence to: Lidong Bing <lidong.bing@shanda.com>.




Figure 1. Multi-agent TIP framework. The planner-agent decomposes queries into subtasks and delegates them to worker-agents, which execute via single-agent TIP and return results. The planner then generates successive subtasks or the final answer.

Among the available tools, search APIs are particularly crucial, allowing LLMs to access external knowledge far beyond their parametric memory. Current implementations typically employ a single-agent TIP paradigm (Dong et al., 2026; Jin et al., 2025), where one LLM instance handles the entire planning and tool-calling process. However, this approach faces two key limitations: 1. **High token cost**: tool responses (e.g., web search results) often consume excessive tokens, making long-horizon multi-turn TIP infeasible when the LLM’s context window is limited; 2. **Contextual distraction**: tool responses often contain irregular patterns (e.g., URLs) or information irrelevant to the query, which can distract the LLM’s attention and impair its ability to plan high-quality subsequent actions.

These limitations motivate the **multi-agent TIP** paradigm (Hu et al., 2025; Du et al., 2024), which amortizes the TIP process across multiple agent roles: a planner-agent that decomposes queries into subtasks, and worker-agents that solve these subtasks via single-agent TIP (Figure 1). This design alleviates the high-token-cost issue by decomposing the full rollout into shorter, manageable segments. Moreover, tool-response tokens are localized within each worker’s context and do not affect the planner’s subsequent reasoning. Importantly, these agent roles can be instantiated within a *single* LLM by managing role-specific system prompts, avoiding the need for multiple model deployments. While multi-agent TIP offers clear benefits in saving token costs and mitigating contextual distractions, training such systems

introduces new challenges, particularly when each agent operates on separate model instances, which incurs infrastructure overhead due to uneven workloads and increased parameter consumption. This naturally raises the following research questions: 1. How can we perform multi-agent RL training effectively using a single LLM? 2. How should credit assignment be handled in multi-turn planner-worker interactions? 3. Can the planner- and worker-agent roles coexist and be trained within a single LLM instance?

In this paper, we provide affirmative answers to these questions. We propose Multi-Agent Tool-Integrated Policy Optimization (**MATPO**), a principled RL algorithm for the **multi-agent-in-one-model** setting, where a single LLM serves as both planner and worker through role-specific system prompts. MATPO derives the policy gradient for nested planner-worker rollouts and introduces a $G \times (T+1)$ group-relative normalization that shares the outcome reward across planner and worker sub-rollouts, building upon recent advances in RL post-training (Shao et al., 2024; Noukhovitch et al., 2025; Zhong et al., 2025). It can be implemented on top of existing single-agent RL frameworks (e.g., verL¹) with minimal modification. In practice, MATPO exhibits more stable learning curves than single-agent approaches.

However, in long-horizon planning, producing a correct final answer depends on a sequence of valid intermediate subtasks, whereas MATPO relies on a sparse outcome reward that assigns a single signal to the whole rollout. To provide turn-level supervision, we introduce **MATPO-PR**, which integrates process rewards into the MATPO framework. Inspired by recent advances in process reward models (Lightman et al., 2024; Wang et al., 2024; Khalifa et al., 2025) and LLM-as-judge verification (Li et al., 2025a), MATPO-PR employs an LLM judge to score each intermediate planner-worker interaction and constructs discounted process returns that credit effective intermediate actions, encouraging the agent to reach correct solutions earlier and more reliably.

Contributions. 1. We present MATPO, a principled RL post-training method for the **multi-agent-in-one-model** setting, with a policy-gradient derivation for nested planner-worker rollouts and a $G \times (T+1)$ group-relative normalization across planner and worker sub-rollouts. 2. We introduce MATPO-PR, a process-reward variant that provides turn-level supervision across planner-worker interactions, enabling temporal credit assignment in long-horizon rollouts. 3. MATPO and MATPO-PR are implemented on top of standard single-agent RL frameworks via a nested-rollout and packed-batch design, with minimal modification to the training loop. 4. Experiments on three deep research benchmarks show consistent improvements over single-agent baselines, with MATPO-PR further improving over MATPO.

¹<https://github.com/volcengine/verl>

2. Related Work

2.1. Tool-Integrated Agent Frameworks

Tool-integrated planning (TIP) has emerged as a crucial paradigm for LLMs to tackle complex, knowledge-intensive tasks through iterative reasoning with external tools (Zhao et al., 2023; Li et al., 2024; Xu & Peng, 2025; OpenAI, 2025). Representative single-agent approaches include function-calling LLMs (Yang et al., 2025; Qin et al., 2025), ReAct-style agents (Yao et al., 2023; Li et al., 2025d;b; Tao et al., 2025), structured workflow agents (Team et al., 2025), and web agents (Zhou et al., 2024a). However, single-agent TIP faces fundamental challenges: 1. the LLM’s context window is quickly saturated by lengthy tool responses, hindering scalability to deeper reasoning chains (Zhang et al., 2025; Jin et al., 2024; Xiao et al., 2024); 2. noisy tool outputs can disrupt reasoning and induce cascading errors (Zhou et al., 2024c). To mitigate these issues, recent work explores multi-agent frameworks (Hu et al., 2025; MiroMind, 2025b; Du et al., 2024; Motwani et al., 2025; ?) where planner- and worker-agents collaborate: the planner decomposes tasks and delegates subtasks to workers, containing noisy outputs within worker contexts. Training methodologies for multi-agent LLM systems have also been studied recently, including M-GRPO (?) and ReMA (?). Liu et al. (2025a) introduces training for multi-turn multi-agent zero-sum games, but does not address the TIP setting.

2.2. Tool-integrated Agentic Reinforcement Learning

Reinforcement learning with verifiable rewards (RLVR) has proven effective for single-agent TIP (Shao et al., 2024; Jin et al., 2025; ?; MiroMind, 2025a; Nguyen et al., 2025), with recent advances introducing more efficient algorithms such as asynchronous RLHF (Noukhovitch et al., 2025) and token-level optimization (Zhong et al., 2025). Trajectory filtering has been explored for math problem solving with code (Li et al., 2025e; Xue et al., 2025; Feng et al., 2025) and GUI tasks (Dong et al., 2026). Another line of work combines SFT or DPO (Rafailov et al., 2023) on cold-start trajectories with RLVR (Li et al., 2025b; Tao et al., 2025; Wei et al., 2025; Ouyang et al., 2025; Li et al., 2025c; MiroMind, 2025a). A key challenge in multi-turn RL is credit assignment across interaction steps. Hierarchical approaches like ArCHer (Zhou et al., 2024b) address this through actor-critic frameworks, while process reward models (Lightman et al., 2024; Wang et al., 2024; Khalifa et al., 2025) provide step-level supervision for reasoning tasks. LLM-as-judge methods (Li et al., 2025a; Liu et al., 2025b) offer scalable alternatives to human annotation. Self-improvement through multi-agent interaction (Qu et al., 2024; Chen et al., 2025b) has also shown promise. While these methods show gains in single-agent settings, principled extensions to the multi-agent-in-one-model regime remain underexplored.

3. Problem Setup

3.1. Single-Agent Multi-Turn Reinforcement Learning

We first recap single-agent multi-turn RL before extending to the multi-agent setting. Let $\pi_\theta(\cdot|\cdot)$ be an LLM parameterized by θ . For each query $q \sim \mathcal{D}$, the agent aims to generate the correct answer via multi-turn TIP (Figure 2).

Recent works (Dong et al., 2025; Qian et al., 2025) have shown that RLVR is a promising approach for optimizing LLMs’ ability to perform the multi-turn TIP process. Suppose $r(\cdot)$ is a reward function that assigns 1 to correct answers and 0 to incorrect ones. The RL objective of the single-agent TIP system is

$$\min_{\theta} J(\pi_{\theta}) \triangleq \mathbb{E}_{q \sim \mathcal{D}, \tau \sim \pi_{\theta}} [r(\tau)], \quad \tau \triangleq [a_1, s_1, \dots, a_T],$$

$$a_t \sim \pi_{\theta}(\cdot | [p_{\text{sys}}, q, a_1, s_1, \dots, s_{t-1}]), \quad s_t \sim \text{Tool}(a_t).$$

Specifically, p_{sys} is the system prompt defining the agent role and tool schema, a_t is the LLM-generated action at turn t including planning and tool-call blocks, $\text{Tool}(\cdot|a_t)$ is the invoked tool conditioned on a_t , s_t is its response, and τ denotes the complete TIP rollout trajectory.

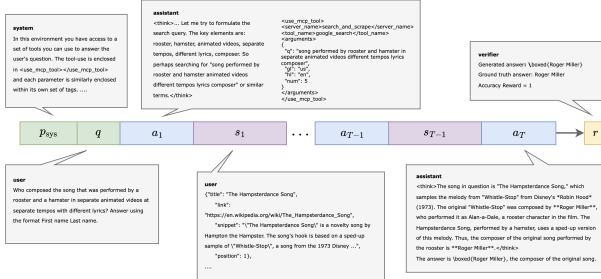


Figure 2. Visualization of a single-agent multi-turn TIP rollout. The LLM solves a query through iterative planning and tool-use. At each step, it plans a tool call, executes it with the parsed parameters, and uses the tool response to decide the next move, continuing until it is confident enough to produce a final answer.

3.2. Single-Agent Group Relative Policy Optimization

GRPO (Shao et al., 2024) is one of the most effective and common methods to minimize $J(\pi_{\theta})$. Each rollout includes LLM-generated tokens $a_1, \dots, a_T$ (blue in Figure 2) and tool responses $s_1, \dots, s_T$ (purple). Since tool responses are not generated by π_{θ} , they do not contribute to the policy gradient. The GRPO objective masks out tool-response tokens as follows:

$$J_{\text{single}}(\pi_{\theta}) \triangleq \mathbb{E}_{\{\tau_i\}_{i=1}^G \sim \pi_{\theta_{\text{old}}}} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{\sum_{t=1}^{T_i} |a_t^i|} \sum_{t=1}^{T_i} R_{i,t}^{\text{clip}} \right],$$

$$R_{i,t}^{\text{clip}} \triangleq \min(R_{i,t}(\theta) \hat{A}_{i,t}, \text{clip}(R_{i,t}(\theta), 1 - \varepsilon, 1 + \varepsilon) \hat{A}_{i,t}),$$

$$R_{i,t}(\theta) \triangleq \frac{\pi_{\theta}(a_t^i | [p_{\text{sys}}, q, a_1^i, s_1^i, \dots, s_{t-1}^i])}{\pi_{\theta_{\text{old}}}(a_t^i | [p_{\text{sys}}, q, a_1^i, s_1^i, \dots, s_{t-1}^i])},$$

$$\hat{A}_{i,t} \triangleq (r(\tau_i) - \text{mean}(\{r(\tau_i)\}_{i=1}^G)) / \text{std}(\{r(\tau_i)\}_{i=1}^G),$$

where $\pi_{\theta_{\text{old}}}$ is a periodically updated snapshot of π_{θ} , G is the group size, $\tau_i \triangleq [a_1^i, s_1^i, \dots, a_{T_i}^i]$ is the i -th rollout with T_i turns, $R_{i,t}(\theta)$ is the likelihood ratio between π_{θ} and $\pi_{\theta_{\text{old}}}$, $\hat{A}_{i,t}$ is the group-relative normalized reward, and $\text{clip}(\cdot, 1 - \varepsilon, 1 + \varepsilon)$ restricts values to $[1 - \varepsilon, 1 + \varepsilon]$.

3.3. Multi-Agent Multi-Turn Reinforcement Learning

As mentioned in the “introduction” section, multi-agent multi-turn TIP frameworks are designed to overcome the context length bottleneck and noisy tool-token issues present in single-agent multi-turn TIP. For clarity and without loss of generality, this paper considers a multi-agent framework with one planner-agent and one worker-agent. A multi-agent multi-turn TIP rollout is visualized in Figure 3. Specifically, q denotes the user query, and τ represents the entire multi-turn TIP rollout for handling it. p_{planner} is the system prompt specifying the role of the planner agent. At each turn t , the planner generates an action a_t containing a thinking block and either a subtask or the final answer, and receives a response s_t parsed from the worker agent’s output. The planner proceeds for T turns in total. Each subtask query $q_{\text{sub},t}$ parsed from a_t is handled by a worker-agent rollout τ^t , guided by the system prompt p_{worker} . Within τ^t , the worker produces actions a_i^t (each including a thinking block and either a tool call or a final sub-answer) and receives tool responses s_i^t . Finally, r denotes the accuracy reward for the final planner answer a_T .




Figure 3. Visualization of a multi-agent multi-turn TIP rollout. The planner assigns subtasks to the worker, which completes them via multi-turn single-agent TIP iterations.

Each multi-agent rollout consists of T single-agent rollouts: one from the planner and $(T - 1)$ from workers. Specifically:

$$\tau \triangleq [a_1, \tau^1, s_1, \dots, a_{T-1}, \tau^{T-1}, s_{T-1}, a_T] \sim (\pi_{\theta}, \text{Tool}),$$

$$\tau^t \triangleq [a_1^t, s_1^t, \dots, s_{T_t-1}^t, a_{T_t}^t], \quad s_t^t \sim \text{Parse}(a_{T_t}^t), \quad s_i^t \sim \text{Tool}(a_i^t),$$

where $\text{Parse}(a_{T_t}^t)$ is the worker-agent’s response to the t -th subtask parsed from the final content in the worker-agent rollout, and $\text{Tool}(a_i^t)$ is the tool-response based on the parameters parsed from action a_i^t from the worker-agent.

Given a reward function $r(\cdot)$ that assigns 1 to correct answers and 0 to incorrect ones, the objective of multi-agent multi-turn RL can be formalized as:

$$\begin{aligned} \min_{\theta} J_{\text{multi}}(\pi_{\theta}) &\triangleq \mathbb{E}_{q \sim \mathcal{D}, \tau \sim (\pi_{\theta}, \text{Tool})} [r(\tau)], \\ a_t &\sim \pi_{\theta}(\cdot | [p_{\text{planner}}, q, a_1, s_1, \dots, s_{t-1}]), \\ a_j^t &\sim \pi_{\theta}(\cdot | [p_{\text{worker}}, q_{\text{sub},t}, a_1^t, s_1^t, \dots, s_{j-1}^t]), \\ s_t &\sim \text{Parse}(a_{T_t}^t), q_{\text{sub},t} \sim \text{Parse}(a_t), s_j^t \sim \text{Tool}(a_j^t). \end{aligned}$$

Notice that in $J_{\text{multi}}(\pi_{\theta})$, a single LLM π_{θ} is deployed to serve as both the planner-agent and the worker-agent, distinguished only by different system prompts p_{planner} and p_{worker} . In this paper, we refer to this deployment configuration as **multi-agent-in-one-model**.

An alternative configuration is to deploy separate models for the planner-agent and worker-agents, which we refer to as multi-agent-multi-model. The multi-agent multi-turn RL objective can be directly generalized to this configuration. Let the planner-agent be parameterized by π_{θ} and K worker-agents be parameterized by $\pi_{\phi_1}, \dots, \pi_{\phi_K}$. The resulting multi-agent-multi-model objective is

$$\begin{aligned} J_{\text{multi}}(\pi_{\theta}, \{\pi_{\phi_k}\}_{k \in [K]}) &\triangleq \mathbb{E}_{q \sim \mathcal{D}, \tau \sim (\pi_{\theta}, \{\pi_{\phi_k}\}_{k \in [K]}, \text{Tool})} [r(\tau)], \\ a_t &\sim \pi_{\theta}(\cdot | [p_{\text{planner}}, q, a_1, s_1, \dots, s_{t-1}]), \\ a_j^t &\sim \pi_{\phi_k}(\cdot | [p_{\text{worker}}, q_{\text{sub},t}, a_1^t, s_1^t, \dots, s_{j-1}^t]), k \in [K], \\ s_t &\sim \text{Parse}(a_{T_t}^t), (q_{\text{sub},t}, k) \sim \text{Parse}(a_t), s_j^t \sim \text{Tool}(a_j^t). \end{aligned}$$

In this paper, we focus on exploring RL training under the multi-agent-in-one-model setting, as it offers several advantages over the multi-agent-multi-model setting: 1. The multi-agent-multi-model setting requires $(K + 1)$ LLM rollout engines and additional RL infrastructure optimization. In contrast, the multi-agent-in-one-model framework uses only one single LLM rollout engine and remains compatible with off-the-shelf RL frameworks; 2. The multi-agent-in-one-model setting allows us to investigate whether RL training can benefit the model when it is exposed to experience from multiple agent roles.

4. Methodology

4.1. Multi-Agent Tool-Integrated Policy Optimization

A key challenge in the multi-agent-in-one-model setting is **credit assignment**: how should the planner-agent rollout τ^0 and the worker-agent rollouts τ^t share responsibility for the final accuracy of the full multi-turn TIP rollout τ ? The planner-agent’s final answer is directly verifiable, whereas worker-agent rollouts address unverifiable subtasks, making it essential to assess their contribution to the planner’s final answer.

The core philosophy of MATPO is that, although the planner and worker have different roles, they exist within the *same rollout lifecycle*. Because their combined actions determine the success of a task, they are optimized against the

same terminal reward, ensuring their policies are perfectly aligned.

We now derive the GRPO counterpart in the multi-agent-in-one-model setting to optimize $J_{\text{multi}}(\pi_{\theta})$. The policy gradient $\nabla_{\theta} J_{\text{multi}}(\pi_{\theta})$ equals

$$\begin{aligned} \nabla_{\theta} J_{\text{multi}}(\pi_{\theta}) &= \nabla_{\theta} \mathbb{E}_{q \sim \mathcal{D}, \tau \sim (\pi_{\theta}, \text{Tool})} [r(\tau)] \\ &= \mathbb{E}_{q \sim \mathcal{D}, \tau \sim (\pi_{\theta}, \text{Tool})} [r(\tau) \nabla_{\theta} \log \mathbb{P}_{\theta}(\tau)], \end{aligned}$$

where $r(\tau)$ is the accuracy reward for rollout τ and $\mathbb{P}_{\theta}(\tau)$ is the generation probability. Let $c_t^p \triangleq [p_{\text{planner}}, q, a_1, s_1, \dots, s_{t-1}]$ and $c_j^{t,w} \triangleq [p_{\text{worker}}, q_{\text{sub},t}, a_1^t, s_1^t, \dots, s_{j-1}^t]$ denote the planner and worker contexts, respectively. By the chain rule, the trajectory probability factorizes along planner-worker boundaries as

$$\begin{aligned} \mathbb{P}_{\theta}(\tau) &= \prod_{t=1}^T \pi_{\theta}(a_t | c_t^p) \cdot \prod_{t=1}^{T-1} \mathbb{P}_{\theta}(\tau^t | a_t), \\ \mathbb{P}_{\theta}(\tau^t | a_t) &= \prod_{j=1}^{T_t} \pi_{\theta}(a_j^t | c_j^{t,w}) \cdot \prod_{j=1}^{T_t-1} \mathbb{P}_{\text{Tool}}(s_j^t | a_j^t). \end{aligned}$$

As the tool-responses are not generated by the LLM π_{θ} , it holds that $\nabla_{\theta} \mathbb{P}_{\text{Tool}}(s_j^t | a_j^t) = 0$, so taking the log-gradient yields

$$\begin{aligned} \nabla_{\theta} \log \mathbb{P}_{\theta}(\tau) &= \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t | c_t^p) \\ &\quad + \sum_{t=1}^{T-1} \sum_{j=1}^{T_t} \nabla_{\theta} \log \pi_{\theta}(a_j^t | c_j^{t,w}). \end{aligned}$$

A step-by-step derivation is provided in Appendix A.1. where $\tau^0 \triangleq [p_{\text{planner}}, q, a_1, s_1, \dots, s_{T-1}, a_T]$ is the planner rollout and τ^t is the t -th worker rollout.

Following standard GRPO derivation, the MATPO objective can be written as:

$$\begin{aligned} J_{\text{MATPO}}(\pi_{\theta}) &\triangleq \mathbb{E}_{\{\tau_i\} \sim (q \sim \mathcal{D}, \pi_{\theta_{\text{old}}}, \text{Tool})} \left[\frac{1}{G} \sum_{i=1}^G \frac{1}{\sum_{t=0}^{T_i} |\tau_i^t|} \sum_{t=0}^{T_i} R_i^{\text{clip}} \right], \\ R_i^{\text{clip}} &\triangleq \min(R_{i,t}, \hat{A}_{i,t}, \text{clip}(R_{i,t}, 1 - \epsilon, 1 + \epsilon) \hat{A}_{i,t}), \\ \hat{A}_{i,t} &\triangleq (r(\tau_i) - \text{mean}(\{r(\tau_i)\}_{i=1}^G)) / \text{std}(\{r(\tau_i)\}_{i=1}^G), \end{aligned}$$

where τ_i is the full rollout for query q with T_i subtasks, τ_i^0 is the planner rollout, τ_i^t ($t > 0$) is the t -th worker rollout, and $\hat{A}_{i,t}$ is the group-relative normalized reward. $R_{i,t}$ is the likelihood ratio:

$$R_{i,t} \triangleq \begin{cases} \sum_{j=1}^{T_i} \frac{\pi_{\theta_{\text{old}}}(a_j^t | [p_{\text{planner}}, q, a_1, s_1, \dots, s_{j-1}])}{\pi_{\theta}(a_j^t | [p_{\text{planner}}, q, a_1, s_1, \dots, s_{j-1}])}, & t = 0, \\ \sum_{j=1}^{T_{i,t}} \frac{\pi_{\theta_{\text{old}}}(a_j^t | [p_{\text{worker}}, q_{\text{sub},t}, a_1^t, s_1^t, \dots, s_{j-1}^t])}{\pi_{\theta}(a_j^t | [p_{\text{worker}}, q_{\text{sub},t}, a_1^t, s_1^t, \dots, s_{j-1}^t])}, & t > 0, \end{cases}$$

where $T_{i,t}$ is the tool-calls count in the t -th subtask of τ_i .




Figure 4. Visualization of MATPO-PR rollout. MATPO-PR assigns turn-level advantages by computing a discounted process return R_t from binary LLM-judged process rewards $\{r_i^{\text{proc}}\}_{i=t}^T$, enabling temporal credit assignment across planner-worker interactions instead of uniform rollout-level supervision.

We summarize the key distinctions between single-agent GRPO and MATPO as follows: unlike GRPO, which performs a single worker-agent rollout per update, MATPO executes one planner-agent rollout followed by T worker-agent rollouts. Moreover, while GRPO normalizes rewards across G worker rollouts for credit assignment, MATPO normalizes across $G \times (T + 1)$ rollouts to jointly account for planner and worker contributions.

4.2. MATPO with Process Rewards

A limitation of standard MATPO is that it solely relies on sparse outcome-based rewards $r(\tau_i)$, which are assigned only upon the completion of the full rollout trajectory, leaving the extensive intermediate steps without direct supervision. This sparsity hinders the optimization of intermediate planner-worker interactions that are crucial for final success. To address this issue, we propose MATPO with Process Reward (**MATPO-PR**), which enables temporal credit assignment by attributing future success to specific intermediate planner-worker interactions. Specifically, upon receiving the worker response s_t at turn t , an LLM judge evaluates its information sufficiency for answering q . We assign a binary **process reward** $r_t^{\text{proc}} = 1$ if sufficient, and 0 otherwise. To facilitate temporal credit assignment, the **process return** $R_t \triangleq \sum_{k=t}^T \gamma^{k-t} r_k^{\text{proc}}$ is employed to quantify the value of the t -th planner-worker interaction, where R_t aggregates future process rewards via a discount factor γ . Intuitively, R_t captures the long-term value of the t -th interaction: a higher value implies that this step is beneficial in enabling subsequent high-quality actions and achieving a reliable final answer.

For each rollout group $\{\tau_i\}_{i=1}^G$ associated with each query

q , we denote $R_{i,t}$ as the t -th process return of the i -th rollout. The MATPO-PR objective is modified from $J_{\text{MATPO}}(\theta_\pi)$ by replacing the outcome reward $r(\tau_i)$ with the process return $R_{i,t}$. Specifically, we replace the group-relative normalized reward $\hat{A}_{i,t}$ with $\hat{A}_{i,t}^{\text{proc}} \triangleq (R_{i,t} - \mu)/\sigma$, where μ and σ denote the mean and standard deviation of $\{R_{i,1}, \dots, R_{i,T^i}\}_{i=1}^G$ over the group of all intermediate interactions within G rollouts, where T^i denotes the number of interactions in the i -th rollout τ_i .

Unlike standard MATPO, which assigns a uniform advantage to all tokens within a rollout, MATPO-PR enables fine-grained turn-level credit assignment, effectively incentivizing the proposal and accomplishment of subtasks that are instrumental to the final answer throughout the long-horizon trajectory.

4.3. Implementation Design

MATPO can be implemented on top of existing single-agent RL frameworks (*e.g.*, veRL). Figure 5 illustrates our implementation, which involves two key architectural designs.




Figure 5. Implementation design of MATPO, showing nested rollouts and packed training batches.

Nested Rollout Mechanism. To reuse state-of-the-art rollout engine implementations, we implement multi-agent rollouts via a nested rollout mechanism: when a worker-agent is invoked within a planner rollout, a nested rollout function is launched inside the outer one, and these processes execute asynchronously. For each query, we feed `n.rollout` requests to the rollout engine (*e.g.*, vLLM and SGLang), generating `n.rollout` planner rollouts (purple boxes in

Figure 5), each with associated worker rollouts (orange boxes) for subtasks.

Training with Packed Rollouts. To ensure compatibility with data parallelism protocols, we pack worker-agent rollouts onto their corresponding planner-agent rollouts. Both are converted to data batches. We compute accuracy rewards by verifying final answers, normalize advantages across planner rollouts per query, and broadcast advantages to corresponding worker rollouts. Finally, planner and worker rollouts are concatenated into an augmented batch. We compute log likelihoods using π_θ and $\pi_{\theta_{old}}$, calculate $J_{MATPO}(\pi_\theta)$ while masking system prompts, queries, and tool responses, then update π_θ via standard optimization.

5. Experiments

5.1. Setups

In this work, we focus on the deep search scenario, where a planner-agent and a worker-agent comprise a two-agent system, aiming to find the answer to a given user query based on searching and web scraping². Specifically, we implement our algorithm on top of veRL. The training hyperparameters are provided in the training script released in the GitHub repository. All experiments are conducted with 128 NVIDIA A800 80G GPUs.

Dataset and Base Model. All experiments are conducted on the Qwen3-14B-base model. We train the model with single-agent GRPO, MATPO, or MATPO-PR on a filtered subset of the MuSiQue (Trivedi et al., 2022) dataset, a multi-hop QA dataset. We remove overly difficult queries for which LLMs repeatedly fail to produce valid rollouts. Our models are then tested on GAIA-text (Mialon et al., 2024)³, WebWalkerQA (Wu et al., 2025), and FRAMES (Krishna et al., 2025). All methods in our comparison share the same base model, training data, tool stack, and evaluator, so the comparison isolates the effect of the training algorithm.

Reward Function. In this work, we employ LLM-as-a-judge⁴ to evaluate the accuracy of a model’s answer against the ground-truth answer. The RL reward is set as $\text{reward} = 0.9 * \text{acc} + 0.1 * \text{fmt}$, where acc is a binary value indicating whether the rollout is correct, fmt measures the average correctness of the tool-calls generated by the model. Specifically, for single-agent RL, we define fmt as the success rate of all tool-call attempts parsed from the LLM’s

generated action. For MATPO, we define $\text{fmt} = 0.5 * \text{fmt_p} + 0.5 * \text{fmt_w}$, where fmt_p denotes the successful tool-call rate among a planner-agent rollout, and fmt_w denotes the average successful tool-call rate among all associated worker-agent rollouts. For MATPO-PR, the settings are similar.

Rollout Summary Mechanism. To encourage the agent to generate answers based on the entire rollout trajectory, we implement a final-summary mechanism. At the end of each rollout, we instruct the model to stop further tool calls and produce an answer based on a summary of the full rollout. We then perform an additional round of summarization and append this final summary to the complete rollout trajectory⁵. To avoid exceeding the model’s maximum token length, if a rollout reaches the limit, we remove the latest messages from the trajectory until there is sufficient token budget for the final summary. Worker-agents in MATPO and MATPO-PR are equipped with this summary mechanism. Additional implementation details are provided in the Appendix, including MATPO-PR specifics (Appendix A.2.2) and agent system prompts with tool-call formats (Appendix A.2.1).

5.2. Results

MATPO consistently outperforms single-agent GRPO.

Figure 6 presents the testing accuracy on GAIA-text, WebWalkerQA, and FRAMES across different training steps. MATPO consistently surpasses the single-agent GRPO baseline, underscoring the effectiveness of our approach. Specifically, MATPO achieves 36.89%, 30.29%, and 62.14% on GAIA-text, WebWalkerQA, and FRAMES, respectively, compared to 33.89%, 28.97%, and 57.34% for single-agent GRPO, leading to an average relative improvement of 7.26%. Moreover, MATPO exhibits more stable gains as training progresses. For instance, while the performance of single-agent GRPO drops after step 90 on both GAIA-text and FRAMES, MATPO continues to improve. We attribute this divergence to the vulnerability of single-agent training: agentic RL often suffers catastrophic drops in performance due to unstable environmental feedback (e.g., missing or noisy responses from the Serper API). In contrast, MATPO can invoke additional browsing subtasks, enabling the agent to perform more robust searches and maintain steady progress.

MATPO-PR consistently outperforms MATPO. To highlight the stability introduced by incorporating process rewards, we evaluate multiple settings of γ , uniformly sampled from the interval $[0, 1]$, specifically $\{0.2, 0.5, 0.8\}$. A larger γ usually corresponds to larger process reward values. We first plot the accuracy reward curves on MuSiQue during training in Figure 7, where “MATPO-PR(γ)” is our result. As shown, “MATPO-PR” consistently achieves higher accuracy than both MATPO and the single-agent baseline.

²To avoid potential leakage of datasets hosted on HuggingFace, search results from this site are blocked by default, unless noted.

³GAIA-text is a curated subset of 103 text-only queries drawn from the GAIA dataset (Mialon et al., 2024), a benchmark for general AI assistants. Absolute scores on GAIA-text are therefore not directly comparable to numbers reported on the full GAIA benchmark, which additionally involves image, audio, and file-based tool use.

⁴We implement LLM-as-a-judge based on GPT-4o-mini with instructions shown in the Appendix.

⁵The rollout summary prompt is detailed in the Appendix.




Figure 6. Test accuracy on three benchmarks across different training steps. Models are trained on the MusiQue dataset.




Figure 7. We record the evolution of training accuracy on the training set for different methods. All curves exhibit an upward trend, while MATPO-PR achieves the highest accuracy.

This advantage also persists on validation benchmarks, as shown in Figure 6. Our validation benchmarks include GAIA-text, WebWalkerQA, and FRAMES. From the validation accuracy curves, we observe that MATPO-PR remains almost entirely above the MATPO and single-agent curves, especially in the later stages of training. MATPO-PR achieves 40.90%, 33.09%, and 64.20% on GAIA-text, WebWalkerQA, and FRAMES, respectively, compared to 36.89%, 30.29%, and 62.14% for MATPO, leading to an average relative improvement of 7.81%.

We attribute the advantages of process rewards to two main factors. Firstly, we amplify the reward score (i.e., as reward return) at the step t where an accuracy reward $r_t^{\text{proc}} = 1$ first appears. This encourages the agent to reach high reward values more quickly. This effect can also be observed in Figure 8, where the visualization of process rewards for MATPO-PR exhibits high process reward values at an earlier stage. Secondly, we also strengthen the rewards after the first occurrence of $r_t^{\text{proc}} = 1$, preventing cases where a rollout achieves an accuracy of 1 after calling a certain sub-agent but then produces an incorrect answer in a subsequent step. This mechanism helps maintain a high and continuous proportion of high-reward steps within each rollout.

5.3. Visualization of Process Reward

The core motivation of MATPO-PR is to design a rational process reward, enhancing the execution quality of intermediate steps. To clearly illustrate this effect, we record the process rewards $\{r_{i,1}^{\text{proc}}, \dots, r_{i,T}^{\text{proc}}\}_{i=1}^G$ for each rollout group of testing samples at different training steps. Although these

reward sequences vary in length ($\{T^i\}_{i=1}^G$), they capture the underlying patterns of process reward dynamics. To enable a unified comparison across different rollouts and training steps, we interpolate each rollout’s process reward sequence to a fixed length using nearest-neighbor interpolation, preserving original reward values.

Let the fixed sequence length be denoted as T , and the number of training steps as M . This process yields a $T \times M$ matrix, which provides a clear visualization of process reward evolution, as shown in Figure 8. From the figure, we observe that rollouts generated by our method achieve higher process rewards as training progresses, compared to the baseline MATPO. This trend is evidenced by the increased concentration of high-reward regions in the *upper-right portion of the matrix*. Moreover, as training progresses, rollouts produced by our method consistently exhibit a *higher proportion of steps with high process rewards* on testing benchmarks.

5.4. Ablation Studies and Practical Take-Aways

We conduct ablation studies on the key components of MATPO and MATPO-PR. Figure 9 shows the testing accuracy under different ablation settings. Each curve represents the following settings. (Purple, ▼): MATPO-PR(0.5); (Brown, +): MATPO-PR(0.5) + Diff.; (Pink, ×): MATPO-PR(0.5) + Partial; (Orange, ■): MATPO; (Green, ◆): MATPO without user query recapping; (Red, ▲): MATPO without final summary; (Blue, ●): single-agent GRPO. Higher curves reflect better accuracy.

Final summaries are necessary. Comparing Orange and Red curves in Figure 9, we find that adding a worker-agent summary significantly improves performance. Without a final summary, the planner-agent may be forced to consume the raw final block, which is error-prone: 1. Long worker-agent outputs may end with tool-call blocks instead of useful answers; 2. The `<think>...</think>` blocks from worker-agents can distract the planner-agent’s subsequent actions. The final summary mitigates both issues, leading to a cleaner interface between the planner- and worker-agent.

Recapping the original user query to worker-agent improves the multi-agent RL performance. We find that the




Figure 8. Visualization of process rewards ($\{r_{i,1}^{\text{proc}}, \dots, r_{i,T_i}^{\text{proc}}\}_{i=1}^G$) on the GAIA-text test set, as training progresses.




Figure 9. Test accuracy across different training steps for varying ablation settings. Models are trained on the MusiQue dataset.

context provided to the worker-agent (e.g., the input prompt) plays a crucial role in determining multi-agent RL performance. A comparison between the **Orange** and **Green** curves in Figure 9 illustrates this effect: recapping the original user query in the worker-agent’s system prompt results in a substantial performance gain. We hypothesize that user query recapping provides the worker-agent explicit guidance toward fulfilling the original user query, thereby improving both the stability and quality of its browsing trajectory.

Cumulative rewards outperform differential rewards.

We evaluate an alternative reward aggregation strategy, differential rewards, denoted as MATPO-PR+Diff. Specifically, suppose a rollout has a reward list $[r_{i,1}^{\text{proc}}, \dots, r_{i,T_i}^{\text{proc}}]$, where r_{i,T_i}^{proc} is the final answer’s reward. The gains of the process reward can be defined as $[r_{i,1}^{\text{proc}}, \dots, r_{i,T_i}^{\text{proc}} - r_{i,T_i-1}^{\text{proc}}, \dots, r_{i,T_i}^{\text{proc}} - r_{i,T_i-1}^{\text{proc}}]$, which can be viewed as the reward list instead. We then compute the final process returns $\{R_{i,t}\}_{t=1}^{T_i}$ using the “discounted sum of future process rewards” for this new reward list.

Comparing the **Brown** and **Purple** curves in Figure 9, we find that MATPO-PR+Diff leads to worse performance than MATPO-PR with cumulative rewards. The underlying reason is that this design effectively turns rewards into differences between consecutive sub-task executions. When a rollout executes a number of sub-tasks, the process rewards $\{r_{i,t}^{\text{proc}}\}_{t=1}^{T_i}$ may become similar for some consecutive steps, causing the computed differences to approach zero. As a result, the rewards of intermediate steps become sparse again, undermining the intended benefit.

Full-trajectory process returns are necessary. We evaluate another alternative process return strategy, partial-

trajectory returns, denoted as MATPO-PR+Partial. This setting assumes that the process return for a rollout is primarily required for the steps preceding the first occurrence of an accuracy value of 1 (i.e., the first T^j where $r_{i,T^j}^{\text{proc}} = 1$). The rationale is that once the accuracy reaches 1 (indicating that the question can already be answered correctly), additional process return is no longer necessary for the following steps. Formally, suppose a rollout has a reward list $[r_{i,1}^{\text{proc}}, \dots, r_{i,T_i}^{\text{proc}}]$, and let r_{i,T^j}^{proc} denote the first process reward that equals 1. In this case, the prefix $[r_{i,1}^{\text{proc}}, \dots, r_{i,T^j}^{\text{proc}}]$ requires process return with higher priority. Therefore, we apply the computation of “discounted sum of future process” to $[r_{i,1}^{\text{proc}}, \dots, r_{i,T^j}^{\text{proc}}]$ to obtain the process returns for steps $\{1, \dots, T^j\}$, while keeping the remaining $[r_{i,T^j+1}^{\text{proc}}, \dots, r_{i,T_i}^{\text{proc}}]$ unchanged as the process returns.

Comparing the **Pink** and **Purple** curves in Figure 9, we find that MATPO-PR+Partial performs worse than MATPO-PR with full-trajectory returns. This indicates that even after a process reward reaches 1 at a certain step, it can revert to 0 in subsequent steps. As a result, it is necessary to apply “discounted sum of future process” to the entire reward sequence $[r_{i,1}^{\text{proc}}, \dots, r_{i,T_i}^{\text{proc}}]$ to achieve better performance.

6. Conclusions

This work introduces MATPO, a principled reinforcement learning framework for multi-agent-in-one-model training using role-specific prompts, along with MATPO-PR, an extension that addresses credit assignment via process rewards. Empirical evaluation across three deep research benchmarks demonstrates consistent improvements over single-agent baselines in both accuracy and robustness. More broadly, MATPO offers a new perspective on LLM post-training by

unifying multiple agent roles within a single model optimized against shared objectives. This paradigm enables richer learning signals from diverse tool interactions and role-specific experiences. A important future direction is to investigate scaling laws between role diversity and LLMs’ agentic reasoning and planning capabilities.

Impact Statement

This paper presents work whose goal is to advance the field of post-training techniques for agentic foundation models, specifically in the area of reinforcement learning for multi-agent systems with tool integration. Our proposed methods, MATPO and MATPO-PR, aim to improve the capability of LLMs to coordinate multiple agent roles and effectively utilize external tools for complex reasoning tasks.

The potential positive societal impacts of this work include enabling more capable AI assistants for knowledge-intensive tasks such as research, information synthesis, and complex problem-solving. Improved multi-agent coordination could benefit applications in education, scientific discovery, and productivity tools.

As with advances in LLM capabilities more broadly, there are potential risks associated with more capable AI systems, including the possibility of generating misleading information or being misused for harmful purposes. However, our work focuses on improving coordination and tool use in controlled research settings, and we do not believe it introduces risks beyond those already present in the development of general-purpose language models.

We encourage the research community to continue developing responsible deployment practices and safety measures as multi-agent AI systems become more prevalent.

References

- Chen, G., Zhang, Z., Cong, X., Guo, F., Wu, Y., Lin, Y., Feng, W., and Wang, Y. Learning evolving tools for large language models. In *Proceedings of ICLR*, 2025a.
- Chen, Y., Wang, Y., Zhu, S., Yu, H., Feng, T., Zhang, M., Patwary, M., and You, J. Multi-agent evolve: Llm self-improve through co-evolution. *arXiv preprint arXiv:2510.23595*, 2025b.
- Dong, G., Chen, Y., Li, X., Jin, J., Qian, H., Zhu, Y., Mao, H., Zhou, G., Dou, Z., and Wen, J.-R. Tool-star: Empowering llm-brained multi-tool reasoner via reinforcement learning. *arXiv preprint arXiv:2505.16410*, 2025.
- Dong, G., Mao, H., Ma, K., Bao, L., Chen, Y., Wang, Z., Chen, Z., Du, J., Wang, H., Zhang, F., Zhou, G., Zhu, Y., Wen, J.-R., and Dou, Z. Agentic reinforced policy optimization. In *Proceedings of ICLR*, 2026.
- Du, Y., Li, S., Torralba, A., Tenenbaum, J. B., and Mordatch, I. Improving factuality and reasoning in language models through multiagent debate. In *Proceedings of ICML*, 2024.
- Feng, J., Huang, S., Qu, X., Zhang, G., Qin, Y., Zhong, B., Jiang, C., Chi, J., and Zhong, W. Retool: Reinforcement learning for strategic tool use in llms. *arXiv preprint arXiv:2504.11536*, 2025.
- Hu, M., Zhou, Y., Fan, W., Nie, Y., Xia, B., Sun, T., Ye, Z., Jin, Z., Li, Y., Chen, Q., Zhang, Z., Wang, Y., Ye, Q., Ghanem, B., Luo, P., and Li, G. Owl: Optimized workflow learning for general multi-agent assistance in real-world task automation. In *Proceedings of NeurIPS*, 2025.
- Jin, B., Zeng, H., Yue, Z., Yoon, J., Arik, S., Wang, D., Zamani, H., and Han, J. Search-r1: Training llms to reason and leverage search engines with reinforcement learning. In *Proceedings of COLM*, 2025.
- Jin, H., Han, X., Yang, J., Jiang, Z., Liu, Z., Chang, C.-Y., Chen, H., and Hu, X. Llm maybe longlm: Selfextend llm context window without tuning. In *Proceedings of ICML*, 2024.
- Khalifa, M., Agarwal, R., Logeswaran, L., Kim, J., Peng, H., Lee, M., Lee, H., and Wang, L. Process reward models that think. *arXiv preprint arXiv:2504.16828*, 2025.
- Krishna, S., Krishna, K., Mohananey, A., Schwarcz, S., Stambler, A., Upadhyay, S., and Faruqui, M. Fact, fetch, and reason: A unified evaluation of retrieval-augmented generation. In *Proceedings of NAACL*, 2025.
- Li, D., Jiang, B., Huang, L., Beigi, A., Zhao, C., Tan, Z., Bhattacharjee, A., Jiang, Y., Chen, C., Wu, T., Shu, K., Cheng, L., and Liu, H. From generation to judgment: Opportunities and challenges of LLM-as-a-judge. In *Proceedings of EMNLP*, 2025a.
- Li, K., Zhang, Z., Yin, H., Zhang, L., Ou, L., Wu, J., Yin, W., Li, B., Tao, Z., Wang, X., Shen, W., Zhang, J., Zhang, D., Wu, X., Jiang, Y., Yan, M., Xie, P., Huang, F., and Zhou, J. Websailor: Navigating super-human reasoning for web agent. *arXiv preprint arXiv:2507.02592*, 2025b.
- Li, X. A review of prominent paradigms for llm-based agents: Tool use, planning (including rag), and feedback learning. In *Proceedings of COLING*, 2024.
- Li, X., Zhao, R., Chia, Y. K., Ding, B., Joty, S., Poria, S., and Bing, L. Chain-of-knowledge: Grounding large language models via dynamic knowledge adapting over heterogeneous sources. In *Proceedings of ICLR*, 2024.

- Li, X., Jin, J., Dong, G., Qian, H., Zhu, Y., Wu, Y., Wen, J.-R., and Dou, Z. Webthinker: Empowering large reasoning models with deep research capability. In *Proceedings of NeurIPS*, 2025c.
- Li, X., Xu, W., Zhao, R., Jiao, F., Joty, S., and Bing, L. Can we further elicit reasoning in LLMs? critic-guided planning with retrieval-augmentation for solving challenging tasks. In *Proceedings of ACL*, 2025d.
- Li, X., Zou, H., and Liu, P. Torl: Scaling tool-integrated rl. *arXiv preprint arXiv:2503.23383*, 2025e.
- Lightman, H., Kosaraju, V., Burda, Y., Edwards, H., Baker, B., Lee, T., Leike, J., Schulman, J., Sutskever, I., and Cobbe, K. Let’s verify step by step. In *Proceedings of ICLR*, 2024.
- Liu, B., Guertler, L., Yu, S., Liu, Z., Qi, P., Balcells, D., Liu, M., Tan, C., Shi, W., Lin, M., Lee, W. S., and Jaques, N. Spiral: Self-play on zero-sum games incentivizes reasoning via multi-agent multi-turn reinforcement learning. *arXiv preprint arXiv:2506.24119*, 2025a.
- Liu, S., Liu, H., Liu, J., Xiao, L., Gao, S., Lyu, C., Gu, Y., Zhang, W., Wong, D. F., Zhang, S., and Chen, K. CompassVerifier: A unified and robust verifier for LLMs evaluation and outcome reward. In *Proceedings of EMNLP*, 2025b.
- Mialon, G., Fourrier, C., Swift, C., Wolf, T., LeCun, Y., and Scialom, T. Gaia: a benchmark for general ai assistants. In *Proceedings of ICLR*, 2024.
- MiroMind. Mirorl: An mcp-first reinforcement learning framework for deep research agent. *GitHub*, 2025a.
- MiroMind. Miroflow: An open-source agentic framework for deep research. *GitHub*, 2025b.
- Motwani, S. R., Smith, C., Das, R. J., Rafailov, R., Torr, P., Laptev, I., Pizzati, F., Clark, R., and de Witt, C. S. MALT: Improving reasoning with multi-agent LLM training. In *Proceedings of COLM*, 2025.
- Nguyen, X.-P., Pandit, S., Reddy, R. G., Xu, A., Savarese, S., Xiong, C., and Joty, S. Sfr-deepresearch: Towards effective reinforcement learning for autonomously reasoning single agents. *arXiv preprint arXiv:2509.06283*, 2025.
- Noukhovitch, M., Huang, S., Xhonneux, S., Hosseini, A., Agarwal, R., and Courville, A. Faster, more efficient RLHF through off-policy asynchronous learning. In *Proceedings of ICLR*, 2025.
- OpenAI. Deep research system card. *Technical Report*, 2025.
- Ouyang, J., Yan, R., Luo, Y., Cheng, M., Liu, Q., Liu, Z., Yu, S., and Wang, D. Training powerful llm agents with end-to-end reinforcement learning, 2025.
- Qian, C., Acikgoz, E. C., He, Q., Wang, H., Chen, X., Hakkani-Tür, D., Tur, G., and Ji, H. Toolrl: Reward is all tool learning needs. In *Proceedings of NeurIPS*, 2025.
- Qin, S., Zhu, Y., Mu, L., Zhang, S., and Zhang, X. Meta-tool: Unleash open-world function calling capabilities of general-purpose large language models. In *Proceedings of ACL*, 2025.
- Qu, C., Dai, S., Wei, X., Cai, H., Wang, S., Yin, D., Xu, J., and Wen, J.-r. Tool learning with large language models: a survey. *Frontiers of Computer Science*, 2025.
- Qu, Y., Zhang, T., Garg, N., and Kumar, A. Recursive introspection: Teaching LLM agents how to self-improve. In *Proceedings of ICML*, 2024.
- Rafailov, R., Sharma, A., Mitchell, E., Ermon, S., Manning, C. D., and Finn, C. Direct preference optimization: Your language model is secretly a reward model. In *Proceedings of NeurIPS*, 2023.
- Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., Zhang, H., Zhang, M., Li, Y. K., Wu, Y., and Guo, D. Deepseek-math: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Singh, J., Pandya, Y., Vajreshwari, P., Magazine, R., and Nambi, A. Agentic reasoning and tool integration for LLMs via reinforcement learning. In *First Workshop on Foundations of Reasoning in Language Models*, 2025.
- Tao, Z., Wu, J., Yin, W., Zhang, J., Li, B., Shen, H., Li, K., Zhang, L., Wang, X., Jiang, Y., Xie, P., Huang, F., and Zhou, J. Webshaper: Agentic data synthesizing via information-seeking formalization. *arXiv preprint arXiv:2507.15061*, 2025.
- Team, K., Bai, Y., Bao, Y., Chen, G., Chen, J., Chen, N., Chen, R., Chen, Y., Chen, Y., Chen, Y., and et al. Kimi k2: Open agentic intelligence. *arXiv preprint arXiv:2507.20534*, 2025.
- Trivedi, H., Balasubramanian, N., Khot, T., and Sabharwal, A. Musique: Multihop questions via single-hop question composition. In *Proceedings of TACL*, 2022.
- Wang, P., Li, L., Shao, Z., Xu, R., Dai, D., Li, Y., Chen, D., Wu, Y., and Sui, Z. Math-shepherd: Verify and reinforce LLMs step-by-step without human annotations. In *Proceedings of ACL*, 2024.

- Wei, Z., Yao, W., Liu, Y., Zhang, W., Lu, Q., Qiu, L., Yu, C., Xu, P., Zhang, C., Yin, B., Yun, H., and Li, L. Webagent-r1: Training web agents via end-to-end multi-turn reinforcement learning. In *Proceedings of EMNLP*, 2025.
- Wu, J., Yin, W., Jiang, Y., Wang, Z., Xi, Z., Fang, R., Zhang, L., He, Y., Zhou, D., Xie, P., and Huang, F. Webwalker: Benchmarking llms in web traversal. In *Proceedings of ACL*, 2025.
- Xiao, C., Zhang, P., Han, X., Xiao, G., Lin, Y., Zhang, Z., Liu, Z., and Sun, M. Infilmm: training-free long-context extrapolation for llms with an efficient context memory. In *Proceedings of NeurIPS*, 2024.
- Xu, R. and Peng, J. A comprehensive survey of deep research: Systems, methodologies, and applications. *arXiv preprint arXiv:2506.12594*, 2025.
- Xue, Z., Zheng, L., Liu, Q., Li, Y., Zheng, X., Ma, Z., and An, B. Simpletir: End-to-end reinforcement learning for multi-turn tool-integrated reasoning. *arXiv preprint arXiv:2509.02479*, 2025.
- Yang, A., Li, A., Yang, B., Zhang, B., Hui, B., Zheng, B., Yu, B., Gao, C., Huang, C., Lv, C., and et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., and Cao, Y. React: Synergizing reasoning and acting in language models. In *Proceedings of ICLR*, 2023.
- Zhang, J., Zhu, Y., Sun, M., Luo, Y., Qiao, S., Du, L., Zheng, D., Chen, H., and Zhang, N. Lightthinker: Thinking step-by-step compression. In *Proceedings of EMNLP*, 2025.
- Zhao, R., Li, X., Joty, S., Qin, C., and Bing, L. Verify-and-edit: A knowledge-enhanced chain-of-thought framework. In *Proceedings of ACL*, 2023.
- Zhong, H., Shan, Z., Feng, G., Xiong, W., Cheng, X., Zhao, L., He, D., Bian, J., and Wang, L. DPO meets PPO: Reinforced token optimization for RLHF. In *Proceedings of ICML*, 2025.
- Zhou, S., Xu, F. F., Zhu, H., Zhou, X., Lo, R., Sridhar, A., Cheng, X., Ou, T., Bisk, Y., Fried, D., Alon, U., and Neubig, G. Webarena: A realistic web environment for building autonomous agents. In *Proceedings of ICLR*, 2024a.
- Zhou, Y., Zanette, A., Pan, J., Levine, S., and Kumar, A. ArCher: Training language model agents via hierarchical multi-turn RL. In *Proceedings of ICML*, 2024b.
- Zhou, Z., Tao, R., Zhu, J., Luo, Y., Wang, Z., and Han, B. Can language models perform robust reasoning in chain-of-thought prompting with noisy rationales? In *Proceedings of NeurIPS*, 2024c.

A. Appendix

A.1. Step-by-Step Derivation of the MATPO Policy Gradient

This appendix gives the full derivation behind the compact form presented in the main text. Starting from

$$\begin{aligned}\nabla_{\theta} J_{\text{multi}}(\pi_{\theta}) &= \nabla_{\theta} \mathbb{E}_{q \sim \mathcal{D}, \tau \sim (\pi_{\theta}, \text{Tool})} [r(\tau)] \\ &= \mathbb{E}_{q \sim \mathcal{D}, \tau \sim (\pi_{\theta}, \text{Tool})} [r(\tau) \nabla_{\theta} \log \mathbb{P}_{\theta}(\tau)],\end{aligned}$$

we expand the trajectory probability along the planner-worker sequence. The planner emits T actions $a_1, \dots, a_T$ interleaved with $T - 1$ worker sub-rollouts $\tau^1, \dots, \tau^{T-1}$ and their parsed responses $s_1, \dots, s_{T-1}$, giving

$$\begin{aligned}\mathbb{P}_{\theta}(\tau) &\triangleq \mathbb{P}_{\theta}([p_{\text{planner}}, q, a_1, \tau^1, s_1, \dots, \tau^{T-1}, s_{T-1}, a_T]) \\ &= \pi_{\theta}(a_1 | [p_{\text{planner}}, q]) \mathbb{P}_{\theta}(\tau^1 | a_1) \cdots \mathbb{P}_{\theta}(\tau^{T-1} | a_{T-1}) \\ &\quad \cdot \pi_{\theta}(a_T | [p_{\text{planner}}, q, a_1, \dots, s_{T-1}]).\end{aligned}$$

Each worker sub-rollout $\tau^t = [a_1^t, s_1^t, \dots, s_{T_t-1}^t, a_{T_t}^t]$ contains T_t worker actions interleaved with $T_t - 1$ tool responses, giving

$$\begin{aligned}\mathbb{P}_{\theta}(\tau^t | a_t) &\triangleq \mathbb{P}_{\theta}([p_{\text{worker}}, q_{\text{sub},t}, a_1^t, s_1^t, \dots, s_{T_t-1}^t, a_{T_t}^t]) \\ &= \pi_{\theta}(a_1^t | [p_{\text{worker}}, q_{\text{sub},t}]) \mathbb{P}_{\text{Tool}}(s_1^t | a_1^t) \\ &\quad \cdot \pi_{\theta}(a_2^t | [p_{\text{worker}}, q_{\text{sub},t}, a_1^t, s_1^t]) \cdots \mathbb{P}_{\text{Tool}}(s_{T_t-1}^t | a_{T_t-1}^t) \\ &\quad \cdot \pi_{\theta}(a_{T_t}^t | [p_{\text{worker}}, q_{\text{sub},t}, a_1^t, \dots, s_{T_t-1}^t]).\end{aligned}$$

Since tool responses are not generated by π_{θ} , we have $\nabla_{\theta} \log \mathbb{P}_{\text{Tool}}(s_j^t | a_j^t) = 0$. Taking the log-gradient of $\mathbb{P}_{\theta}(\tau)$ then expanding each factor yields

$$\begin{aligned}\nabla_{\theta} \log \mathbb{P}_{\theta}(\tau) &= \nabla_{\theta} \left(\log \pi_{\theta}(a_1 | [p_{\text{planner}}, q]) + \log \mathbb{P}_{\theta}(\tau^1 | a_1) + \cdots \right. \\ &\quad \left. + \log \mathbb{P}_{\theta}(\tau^{T-1} | a_{T-1}) + \log \pi_{\theta}(a_T | [p_{\text{planner}}, q, a_1, \dots, s_{T-1}]) \right) \\ &= \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t | [p_{\text{planner}}, q, a_1, s_1, \dots, s_{t-1}]) \\ &\quad + \sum_{t=1}^{T-1} \sum_{j=1}^{T_t} \nabla_{\theta} \log \pi_{\theta}(a_j^t | [p_{\text{worker}}, q_{\text{sub},t}, a_1^t, s_1^t, \dots, s_{j-1}^t]) \\ &= \sum_{t=1}^T \frac{\nabla_{\theta} \pi_{\theta}(a_t | [p_{\text{planner}}, q, a_1, s_1, \dots, s_{t-1}])}{\pi_{\theta}(a_t | [p_{\text{planner}}, q, a_1, s_1, \dots, s_{t-1}])} \\ &\quad + \sum_{t=1}^{T-1} \sum_{j=1}^{T_t} \frac{\nabla_{\theta} \pi_{\theta}(a_j^t | [p_{\text{worker}}, q_{\text{sub},t}, a_1^t, s_1^t, \dots, s_{j-1}^t])}{\pi_{\theta}(a_j^t | [p_{\text{worker}}, q_{\text{sub},t}, a_1^t, s_1^t, \dots, s_{j-1}^t])},\end{aligned}$$

matching the compact form given in the main text. Substituting back into $\nabla_{\theta} J_{\text{multi}}$ and applying standard GRPO clipping with group-relative normalization yields the J_{MATPO} objective.

A.2. Implementation Details

A.2.1. AGENT SYSTEM PROMPT AND TOOL-CALL FORMAT

We adopt an XML format to parse tool calls from both planner and worker agents. The planner-agent’s system prompt specifies the tool schema for invoking the worker-agent, while the worker-agent’s system prompt defines the schema

for Google’s Serper API (used for search and scraping). After each tool call, the tool’s response is wrapped as a “user message” and appended to the agent’s rollout trajectory. To help the worker agent execute the user’s original query from the planner agent, we include a recap of the query in the worker agent’s system prompt, a technique we refer to as *user query recapping*. The detailed system prompts and tool schemas are provided in Sec. A.4.

A.2.2. IMPLEMENTATION DETAILS FOR MATPO-PR

The process reward is computed by requiring the main agent to produce a summary based on the information available up to a specific step, i.e., immediately after completing a sub-task. After the process reward is obtained, the summary operation and its generated content are removed from the main agent’s trajectory. As a result, this auxiliary summarization step does not affect the agent’s trajectory. This design preserves the original workflow of MATPO while augmenting it with additional process-level supervision.

We place the summary operation after all sub-agents complete their tasks, as the main agent may invoke multiple sub-agents to accomplish different sub-tasks within a single step. In this case, the summary operation is executed only after all sub-agents have finished their respective sub-tasks. Consequently, a single, well-defined reward value is assigned at this step to serve as the process reward.

Similar to MATPO, the format reward for the main agent is defined as the average of the format reward computed from the main agent’s rollout and the mean format reward across all sub-agents’ rollouts, as $\text{fmt} = 0.5 * \text{fmt_p} + 0.5 * \text{fmt_w}$. Conversely, the format reward for each sub-agent is defined as the average of the format reward from its own rollout (fmt_wi) and the format reward from the main agent’s rollout (fmt_p), as $\text{fmt} = 0.5 * \text{fmt_wi} + 0.5 * \text{fmt_p}$.

We first allocate a dedicated storage to record the intermediate summaries generated after invoking the sub-agents. These summaries are then evaluated by an LLM-as-judge to produce corresponding scores, which are used to derive the process rewards. During the evaluation stage, each summary is sent to the LLM independently for scoring.

A.3. Analysis of MATPO-PR

A.3.1. ANALYSIS OF THE NORMALIZATION STRATEGY
When performing multi-agent GRPO, it is necessary to normalize the rewards in the rollout group $\{\tau_i\}_{i=1}^G$ to stabilize RL training. However, unlike the previous MATPO setting, the introduction of process rewards means that each rollout now contains multiple reward values. The normalization strategy must explicitly account for all process rewards.

We consider two possible approaches to address this. 1. During training, suppose that for each task there are G




Figure 10. We record the mean and standard deviation of rollout rewards ($\{r_{i,1}^{\text{proc}}, \dots, r_{i,T^i}^{\text{proc}}\}_{i=1}^G$) before normalization (a and b), and compute the mean, maximum, minimum, and standard deviation after normalization (c, d, e, and f). As shown, using *rollout-level* rewards to estimate normalization parameters can produce extreme normalized values that destabilize training, whereas *step-level* normalization yields more stable and well-behaved reward distributions. These statistics are reported for MuSiQue under the “MATPO-PR ($\gamma = 0.5$)” setting at training step 60.

rollouts (as $\{\tau_i\}_{i=1}^G$), and the i -th rollout contains $T^i - 1$ sub-tasks. In this case, the reward sequence of the i -th rollout has length T^i . We take the average of these T^i reward values as the “rollout-level reward” for the i -th rollout ($\{r_{i,1}^{\text{proc}}, \dots, r_{i,T^i}^{\text{proc}}\}$), obtaining μ_i . Then we compute the mean and standard deviation (μ and σ) over $\{\mu_i\}_{i=1}^G$ to perform normalization in GRPO. 2. We first collect the reward lists from all G rollouts, as $\{r_{1,1}^{\text{proc}}, \dots, r_{1,T^1}^{\text{proc}}, \dots, r_{G,1}^{\text{proc}}, \dots, r_{G,T^G}^{\text{proc}}\}$. We then compute the mean and standard deviation over all reward values in this list jointly, and use them to perform normalization in GRPO. The former uses **rollout-level** normalization parameters, whereas the latter adopts more fine-grained, **step-level** normalization parameters.

However, we demonstrate that the first approach is problematic. With the introduction of process rewards, rollouts within a group can have different reward sequences while sharing the same mean value. In such cases, the standard

Pattern Type	Previous released MATPO rollouts	Reproduced rollouts using released MATPO model	Rollouts of new trained MATPO
<i>peopleAlsoAsk</i>	400/1786=22.39%	2251/2811=80.08%	1312/1731=75.79%
<i>relatedSearches</i>	805/1786=45.07%	2382/2811=84.74%	1344/1731=77.64%
<i>answerBox</i>	169/1786=9.46%	896/2811=31.87%	540/1731=31.20%

Table 1. Statistics of the Serper parameters used in Google Search: the proportion of trajectories containing the corresponding returned parameters in Google Search-based executions.

deviation becomes **zero**, leading to extreme values in the normalized rewards, as illustrated in Figure 10. The second normalization method avoids this issue by directly accounting for individual process reward values across all rollouts. For sub-agents, we apply the normalization parameters derived from the main agent.

A.3.2. NOTE ON RESULT VARIABILITY DUE TO EXTERNAL TOOLS

During our experiments, we observed that agent performance (e.g., MATPO) may vary over time due to changes in the external tools and servers they access.

Upon investigation, we identified the primary cause as changes in the Serper API, a critical tool for Google search and web scraping. Specifically, we found that this API has recently undergone parameter changes, particularly in its Google search component.

Specifically, we analyzed the Google search parameters used in the previous MATPO rollouts (run several months ago) and compared them with those used in our current experiments (a period shortly before submission), focusing on three parameters: *peopleAlsoAsk*, *relatedSearches*, and *answerBox*. We found that the proportions of these parameters differed substantially between the two settings. In the current experiments, these parameters are returned at a much higher rate, whereas they appeared only rarely in the previous MATPO runs. As a result, the Google search calls in the current Serper configuration return significantly more information, which may introduce noisy or distracting content and substantially increase input length. Both factors can potentially degrade accuracy. Detailed statistics of these parameters are reported in Table 1.

A.4. Prompts

A.4.1. SYSTEM PROMPT AND TOOL SCHEMA OF THE PLANNER-AGENT

System Prompt:

```
1 In this environment you have
  access to a set of tools you
  can use to answer the user's
  question.
```

```
2
3 You only have access to the tools
  provided below. You can only
  use one tool per message, and
  will receive the result of
  that tool in the user's next
  response. You use tools step-
  by-step to accomplish a given
  task, with each tool-use
  informed by the result of the
  previous tool-use. Today is:
  2025-07-16
4
5 # Tool-Use Formatting Instructions
6
7 Tool-use is formatted using XML-
  style tags. The tool-use is
  enclosed in <use_mcp_tool></
  use_mcp_tool> and each
  parameter is similarly
  enclosed within its own set of
  tags.
8
9 The Model Context Protocol (MCP)
  connects to servers that
  provide additional tools and
  resources to extend your
  capabilities. You can use the
  server's tools via the `
  use_mcp_tool`.
10
11 Description:
12 Request to use a tool provided by
  a MCP server. Each MCP server
  can provide multiple tools
  with different capabilities.
  Tools have defined input
  schemas that specify required
  and optional parameters.
13
14 Parameters:
15 - server_name: (required) The name
  of the MCP server providing
  the tool
16 - tool_name: (required) The name
  of the tool to execute
17 - arguments: (required) A JSON
  object containing the tool's
  input parameters, following
  the tool's input schema,
  quotes within string must be
  properly escaped, ensure it's
  valid JSON
18
19 Usage:
20 <use_mcp_tool>
21 <server_name>server name here</
  server_name>
22 <tool_name>tool name here</
  tool_name>
23 <arguments>
24 {
25   "param1": "value1",
```

```

26  "param2": "value2 "escaped string
27  }
28  </arguments>
29  </use_mcp_tool>
30
31  Important Notes:
32  - Tool-use must be placed **at the
    end** of your response, **top
    -level**, and not nested
    within other tags.
33  - Always adhere to this format for
    the tool use to ensure proper
    parsing and execution.
34
35  String and scalar parameters
    should be specified as is,
    while lists and objects should
    use JSON format. Note that
    spaces for string values are
    not stripped. The output is
    not expected to be valid XML
    and is parsed with regular
    expressions.
36
37
38
39  Here are the functions available
    in JSONSchema format:
40
41  ## Server name: browsing_agent
42  ### Tool name: search_and_browse
43  Description: This tool is an agent
    that performs the subtask of
    searching and browsing the web
    for specific missing
    information and generating the
    desired answer. The subtask
    should be clearly defined,
    include relevant background,
    and focus on factual gaps. It
    does not perform vague or
    speculative subtasks.
44  Args:
45    subtask: the subtask to be
    performed.
46  Returns:
47    the result of the subtask.
48  Input JSON schema: {'properties':
    {'subtask': {'title': 'Subtask',
    'type': 'string'}}, '
    required': ['subtask'], 'title':
    'search_and_browseArguments',
    'type': 'object'}
49
50
51
52  # General Objective
53
54  You accomplish a given task
    iteratively, breaking it down
    into clear steps and working
    through them methodically.

```

```

55
56  ## Task Strategy
57
58  1. Analyze the user's request and
    set clear, achievable sub-
    goals. Prioritize these sub-
    goals in a logical order.
59  2. Start with a concise, numbered,
    step-by-step plan outlining
    how you will solve the task
    before taking any action.
60  3. Work through these sub-goals
    sequentially. After each step,
    adjust your plan as needed.
61  4. Use tools strategically to
    accomplish each sub-goal.
62  5. Revise earlier steps if new
    information emerges.
63
64  ## Tool-Use Guidelines
65
66  1. Each step must involve a single
    tool call, unless the task is
    already solved.
67  2. Before each tool call:
68    - Summarize what is known.
69    - Identify what is missing.
70    - Choose the most relevant tool
    .
71    - Verify all required
    parameters.
72  3. All tool queries must include
    full context.
73  4. Avoid vague queries. Each call
    should retrieve actionable
    information.
74  5. Extract and summarize partial
    information if a tool result
    is incomplete.
75
76  ## Tool-Use Communication Rules
77
78  1. Do not include tool results in
    your response.
79  2. Do not present the final answer
    until the entire task is
    complete.
80  3. Do not mention tool names.
81  4. Do not engage in unnecessary
    back-and-forth.
82  5. Do not use non-existent tools.
83  6. Respond in the same language as
    the user's message.
84  7. If the task does not require
    tool use, answer directly.
85
86
87  # Agent Specific Objective
88
89  You are a task-solving agent that
    uses tools step-by-step to
    answer the user's question.
    Your goal is to provide

```

complete, accurate and well-reasoned answers using additional tools.

A.4.2. SYSTEM PROMPT AND TOOL SCHEMA OF THE WORKER-AGENT

System Prompt:

```

1 In this environment you have
  access to a set of tools you
  can use to answer the user's
  question.
2
3 You only have access to the tools
  provided below. You can only
  use one tool per message, and
  will receive the result of
  that tool in the user's next
  response. You use tools step-
  by-step to accomplish a given
  task, with each tool-use
  informed by the result of the
  previous tool-use. Today is:
  2025-07-08
4
5 # Tool-Use Formatting Instructions
6
7 Tool-use is formatted using XML-
  style tags. The tool-use is
  enclosed in <use_mcp_tool></
  use_mcp_tool> and each
  parameter is similarly
  enclosed within its own set of
  tags.
8
9 The Model Context Protocol (MCP)
  connects to servers that
  provide additional tools and
  resources to extend your
  capabilities. You can use the
  server's tools via the `
  use_mcp_tool`.
10
11 Description:
12 Request to use a tool provided by
  a MCP server. Each MCP server
  can provide multiple tools
  with different capabilities.
  Tools have defined input
  schemas that specify required
  and optional parameters.
13
14 Parameters:
15 - server_name: (required) The name
  of the MCP server providing
  the tool
16 - tool_name: (required) The name
  of the tool to execute
17 - arguments: (required) A JSON
  object containing the tool's
  input parameters, following

```

the tool's input schema,
quotes within string must be
properly escaped, ensure it's
valid JSON

```

18
19 Usage:
20 <use_mcp_tool>
21 <server_name>server name here</
  server_name>
22 <tool_name>tool name here</
  tool_name>
23 <arguments>
24 {
25   "param1": "value1",
26   "param2": "value2 \"escaped
    string\""
27 }
28 </arguments>
29 </use_mcp_tool>
30
31 Important Notes:
32 - Tool-use must be placed at the
  end of your response, top
  -level, and not nested
  within other tags.
33 - Always adhere to this format for
  the tool use to ensure proper
  parsing and execution.
34
35 String and scalar parameters
  should be specified as is,
  while lists and objects should
  use JSON format. Note that
  spaces for string values are
  not stripped. The output is
  not expected to be valid XML
  and is parsed with regular
  expressions.
36
37
38 Here are the functions available
  in JSONSchema format:
39
40 ## Server name:
  search_and_scrape_webpage
41 ### Tool name: google_search
42 Description: Tool to perform web
  searches via Serper API and
  retrieve rich results. It is
  able to retrieve organic
  search results, people also
  ask, related searches, and
  knowledge graph.
43 Input JSON schema: {'type': '
  object', 'properties': {'q':
  {'type': 'string', '
  description': 'Search query
  string'}, 'gl': {'type': '
  string', 'description': "
  Optional region code for
  search results in ISO 3166-1
  alpha-2 format (e.g., 'us')"},
  'hl': {'type': 'string', '

```

```

description': "Optional
language code for search
results in ISO 639-1 format (e
.g., 'en')", 'location': {'
type': 'string', 'description
': "Optional location for
search results (e.g., 'SoHo,
New York, United States', '
California, United States')",
'num': {'type': 'number', '
description': 'Number of
results to return (default:
10)'}, 'tbs': {'type': 'string
', 'description': "Time-based
search filter ('qdr:h' for
past hour, 'qdr:d' for past
day, 'qdr:w' for past week, '
qdr:m' for past month, 'qdr:y'
for past year)"}, 'page': {'
type': 'number', 'description
': 'Page number of results to
return (default: 1)'}, '
autocorrect': {'type': '
boolean', 'description': '
Whether to autocorrect
spelling in query'}}, '
required': ['q', 'gl', 'hl']]

```

```

44
45 ### Tool name: scrape
46 Description: Tool to scrape a
webpage and retrieve the text
and, optionally, the markdown
content. It will retrieve also
the JSON-LD metadata and the
head metadata.
47 Input JSON schema: {'type': '
object', 'properties': {'url':
{'type': 'string', '
description': 'The URL of the
webpage to scrape.'}, '
includeMarkdown': {'type': '
boolean', 'description': '
Whether to include markdown
content.', 'default': False}},
'required': ['url']]

```

```

48
49
50
51 # General Objective
52
53 You accomplish a given task
iteratively, breaking it down
into clear steps and working
through them methodically.
54
55 ## Task Strategy
56
57 1. Analyze the user's request and
set clear, achievable sub-
goals. Prioritize these sub-
goals in a logical order.
58 2. Start with a concise, numbered,
step-by-step plan outlining

```

```

how you will solve the task
before taking any action.
59 3. Work through these sub-goals
sequentially. After each step,
adjust your plan as needed.
60 4. Use tools strategically to
accomplish each sub-goal.
61 5. Revise earlier steps if new
information emerges.

```

```

62
63 ## Tool-Use Guidelines
64
65 1. Each step must involve a single
tool call, unless the task is
already solved.
66 2. Before each tool call:
67 - Summarize what is known.
68 - Identify what is missing.
69 - Choose the most relevant tool
.
70 - Verify all required
parameters.
71 3. All tool queries must include
full context.
72 4. Avoid vague queries. Each call
should retrieve actionable
information.
73 5. Extract and summarize partial
information if a tool result
is incomplete.

```

```

74
75 ## Tool-Use Communication Rules
76
77 1. Do not include tool results in
your response.
78 2. Do not present the final answer
until the entire task is
complete.
79 3. Do not mention tool names.
80 4. Do not engage in unnecessary
back-and-forth.
81 5. Do not use non-existent tools.
82 6. Respond in the same language as
the user's message.
83 7. If the task does not require
tool use, answer directly.

```

```

84
85 # Agent Specific Objective
86
87 You are a task-solving agent that
uses tools step-by-step to
answer the user's question.
Your goal is to provide
complete, accurate and well-
reasoned answers using
additional tools.

```

A.5. Instruction Prompt for Rollout Summarization

System Prompt:

```

1 [SYSTEM]
2 This is a direct instruction to
  you. This is your final turn.
  You MUST NOT use any tools.
3 Your task is to provide a final,
  structured report summarizing
  all the information you have
  gathered to answer your
  assigned subtask.
4
5 [CONTEXT]
6 The main task was: "{main_query}"
7 Your assigned subtask was: "{
  task_description}"
8 Your assigned subtask was intended
  to help solve the main task.
9
10 [INSTRUCTIONS]
11
12 {failed_instruction}
13
14 Your final response MUST be a
  clear, complete, and
  structured report in markdown
  format.
15 Organize the content into logical
  sections with the following
  headings: '## Conclusion', '##
  Supporting Information', '##
  Observations', and '##
  Contribution to Main Task'.
16
17 - **CRITICAL**: Do NOT include raw
  URLs. Replace any URLs with
  '([link])'.
18 - Your response should only
  contain factual, specific, and
  well-organized information
  based on your previous actions
  .
19 - Do not include speculative
  filler, vague summaries, or
  conversational text.
20
21 Here is an example of the required
  format:
22
23 # Final Response: [Title
  summarizing the subtask]
24
25 ## Conclusion:
26 [A concise summary of your
  findings and the final answer
  for the subtask. Bold key
  information.]
27
28 ## Supporting Information:
29 [Detailed supporting facts, data,
  or quotes you discovered. Use
  bullet points or numbered

```

```

    lists for clarity.]
30 - Source 1: Brief description of
  finding 1.
31 - Source 2: Brief description of
  finding 2.
32
33 ## Observations:
34 [Any additional context,
  confidence level, or notes on
  how the conclusion was reached
  .]
35
36 ## Contribution to Main Task:
37 [Explain how the answer to your
  subtask helps solve the
  overall main task. What are
  the next steps the main agent
  should consider?]

```

A.6. Instruction Prompt for Process Summarization to Compute Process Rewards

System Prompt:

```

1 [SYSTEM]
2 This is a direct instruction to
  you. This is your final turn.
  You MUST NOT use any tools.
3 Your task is to provide a final,
  structured report summarizing
  the entire conversation and
  providing a definitive answer
  to the original task.
4
5 The original task was: "{
  main_query}"
6
7 [INSTRUCTIONS]
8
9 {failed_instruction}
10
11 Your final response MUST be a
  clear, complete, and
  structured report in markdown
  format.
12 Organize the content into logical
  sections with the following
  headings: '## Conclusion', '##
  Supporting Information', and
  '## Observations'.
13
14 - **CRITICAL**: Do NOT include any
  raw URLs.
15 - Your response should only
  contain factual, specific, and
  well-organized information
  based on the conversation.
16 - Do not include speculative
  filler, vague summaries, or
  conversational text.
17 - You should wrap the final answer

```



```

18         to the original task in \\
19         boxed{{}}, in '## Conclusion'.
20
21 Here is an example of the required
22     format:
23
24 # Final Response: [Title
25     summarizing the task]
26
27 ## Conclusion:
28 [A concise summary of your
29     findings and the final answer
30     to the user's question. Bold
31     key information. If the answer
32     is a number or short phrase,
33     provide it directly. You
34     should wrap the final answer
35     to the original task in \\
36     boxed{{}}]
37
38 ## Supporting Information:
39 [Detailed supporting facts, data,
40     or quotes from the
41     conversation that support your
42     conclusion. Use bullet points
43     or numbered lists for clarity
44     .]
45
46 - Source 1: Brief description of
47     finding 1.
48 - Source 2: Brief description of
49     finding 2.
50
51 ## Observations:
52 [Any additional context,
53     confidence level, or notes on
54     how the final answer was
55     derived from the conversation
56     history.]

```

A.7. Instruction Prompt for LLM-as-Judge.

System Prompt:

```

1 Your job is to look at a question,
2 a gold target, and a
3 predicted answer, and then
4 assign a grade of either ["
5 CORRECT", "INCORRECT", "
6 NOT_ATTEMPTED"].
7
8 First, I will give examples of
9 each grade, and then you will
10 grade a new example.
11
12
13 The following are examples of
14 CORRECT predicted answers.
15
16 ```
17 Question: What are the names of
18 Barack Obama's children?
19 Gold target: Malia Obama and Sasha
20 Obama

```

```

21
22 Predicted answer 1: sasha and
23 malia obama
24 Predicted answer 2: most people
25 would say Malia and Sasha, but
26 I'm not sure and would have
27 to double check
28 Predicted answer 3: Barack Obama
29 has two daughters. Their names
30 are Malia Ann and Natasha
31 Marian, but they are commonly
32 referred to as Malia Obama and
33 Sasha Obama. Malia was born
34 on July 4, 1998, and Sasha was
35 born on June 10, 2001.
36
37 ```
38 These predicted answers are all
39 CORRECT because:
40
41 - They fully contain the
42     important information in
43     the gold target.
44 - They do not contain any
45     information that
46     contradicts the gold
47     target.
48 - Only semantic meaning
49     matters; capitalization,
50     punctuation, grammar, and
51     order don't matter.
52 - Hedging and guessing are
53     permissible, provided that
54     the gold target is fully
55     included and the response
56     contains no incorrect
57     information or
58     contradictions.
59
60
61 The following are examples of
62 INCORRECT predicted answers.
63
64 ```
65 Question: What are the names of
66 Barack Obama's children?
67 Gold target: Malia and Sasha
68 Predicted answer 1: Malia.
69 Predicted answer 2: Malia, Sasha,
70 and Susan.
71 Predicted answer 3: Barack Obama
72 does not have any children.
73 Predicted answer 4: I think it's
74 either Malia and Sasha. Or it
75 could be Malia and Jackie. Or
76 it could be Joey and Malia.
77 Predicted answer 4: While I don't
78 know their exact names, I can
79 tell you that Barack Obama has
80 three children.
81 Predicted answer 5: It's possible
82 you may mean Betsy and Olivia.
83 However, you should clarify
84 further details with updated
85 references if necessary. Is
86 that the correct answer?
87 Predicted answer 6: It may be the

```

case that Obama's child is named James. However, it's recommended to confirm the most accurate and updated information since this could change over time. This model may not always reflect the most current information.

31 ```

32 These predicted answers are all INCORRECT because:

- 33 - A factual statement in the answer contradicts the gold target. Incorrect statements that have some hedging (e.g., "it is possible that", "although i'm not sure, i think") are also considered incorrect.

34

35

36 The following are examples of NOT_ATTEMPTED predicted answers.

37 ```

38 Question: What are the names of Barack Obama's children?

39 Gold target: Malia and Sasha

40 Predicted answer 1: I don't know.

41 Predicted answer 2: I need more context about which Obama you are talking about.

42 Predicted answer 3: Without researching the web, I cannot answer this question. However, I can tell you that Barack Obama has two children.

43 Predicted answer 4: Barack Obama has two children. I know that one of them is Malia, but I'm not sure about the other one.

44 ```

45 These predicted answers are all NOT_ATTEMPTED because:

- 46 - The important information in the gold target is not included in the answer.
- 47 - No statements in the answer contradict the gold target.

48

49 Also note the following things:

- 50 - For grading questions where the gold target is a number, the predicted answer needs to be correct to the last significant figure in the gold answer. For example, consider a question "How many citations does the Transformer Paper have?" with gold target "120k".

- 51 - Predicted answers "120k", "124k", and 115k" are all CORRECT.

- 52 - Predicted answers "100k" and "113k" are INCORRECT.

- 53 - Predicted answers "around 100k" and "more than 50k" are considered NOT_ATTEMPTED because they neither confirm nor contradict the gold target.

- 54 - The gold target may contain more information than the question. In such cases, the predicted answer only needs to contain the information that is in the question.

- 55 - For example, consider the question "What episode did Derek and Meredith get legally married in Grey's Anatomy?" with gold target "Season 7, Episode 20: White Wedding". Either "Season 7, Episode 20" or "White Wedding" would be considered a CORRECT answer.

- 56 - Do not punish predicted answers if they omit information that would be clearly inferred from the question.

- 57 - For example, consider the question "What city is OpenAI headquartered in?" and the gold target "San Francisco, California". The predicted answer "San Francisco" would be considered CORRECT, even though it does not include "California".

- 58 - Consider the question "What award did A pretrainer's guide to training data: Measuring the effects of data age, domain coverage, quality, & toxicity win at NAACL '24?", the gold target is "Outstanding Paper Award". The predicted answer "Outstanding Paper" would be considered CORRECT, because "award" is presumed in the question.

- 59 - For the question "What is the height of Jason Wei in meters?", the gold target is "1.73 m". The predicted answer "1.75" would be considered

```

CORRECT, because meters is
specified in the question
.
60 - For the question "What is
the name of Barack Obama's
wife?", the gold target
is "Michelle Obama". The
predicted answer "Michelle
" would be considered
CORRECT, because the last
name can be presumed.
61 - Do not punish for typos in
people's name if it's clearly
the same name.
62 - For example, if the gold
target is "Hyung Won Chung
", you can consider the
following predicted
answers as correct: "
Hyoong Won Choong", "
Hyungwon Chung", or "Hyun
Won Chung".
63
64
65 Here is a new example. Simply
reply with either CORRECT,
INCORRECT, NOT ATTEMPTED. Don'
t apologize or correct
yourself if there was a
mistake; we are just trying to
grade the answer.
66 ```
67 Question: {}
68 Gold target: {}
69 Predicted answer: {}
70 ```
71
72 Grade the predicted answer of this
new question as one of:
73 A: CORRECT
74 B: INCORRECT
75 C: NOT_ATTEMPTED
76
77 Just return the letters "A", "B",
or "C", with no text around it
.

```

A.8. Lemon-Pick MATPO Rollout Trajectory

System Prompt:

```

1 # Query:
2
3 On Cornell Law School website's
legal information institute,
under the fifth section of
federal rules alphabetically,
what word was deleted in the
last amendment to the first
rule in the article that has "
witnesses" in the most titles

```

```

as of 2021? You should follow
the format instruction in the
requestion strictly and wrap
the final answer in \boxed{.
4
5 # Ground-truth answer: \boxed{
inference}
6
7 # Rollout generated by a
intermediate checkpoint of
MATPO:
8
9 system
10
11 In this environment you have
access to a set of tools you
can use to answer the user's
question.
12
13 You only have access to the tools
provided below. You can only
use one tool per message, and
will receive the result of
that tool in the user's next
response. You use tools step-
by-step to accomplish a given
task, with each tool-use
informed by the result of the
previous tool-use. Today is:
2025-07-16
14
15 # Tool-Use Formatting Instructions
16
17 Tool-use is formatted using XML-
style tags. The tool-use is
enclosed in <use_mcp_tool></
use_mcp_tool> and each
parameter is similarly
enclosed within its own set of
tags.
18
19 The Model Context Protocol (MCP)
connects to servers that
provide additional tools and
resources to extend your
capabilities. You can use the
server's tools via the `
use_mcp_tool`.
20
21 Description:
22 Request to use a tool provided by
a MCP server. Each MCP server
can provide multiple tools
with different capabilities.
Tools have defined input
schemas that specify required
and optional parameters.
23
24 Parameters:
25 - server_name: (required) The name
of the MCP server providing
the tool
26 - tool_name: (required) The name

```

```

    of the tool to execute
27 - arguments: (required) A JSON
    object containing the tool's
    input parameters, following
    the tool's input schema,
    quotes within string must be
    properly escaped, ensure it's
    valid JSON
28
29 Usage:
30 <use_mcp_tool>
31 <server_name>server name here</
    server_name>
32 <tool_name>tool name here</
    tool_name>
33 <arguments>
34 {
35   "param1": "value1",
36   "param2": "value2"escaped string
    ""
37 }
38 </arguments>
39 </use_mcp_tool>
40
41 Important Notes:
42 - Tool-use must be placed **at the
    end** of your response, **top
    -level**, and not nested
    within other tags.
43 - Always adhere to this format for
    the tool use to ensure proper
    parsing and execution.
44
45 String and scalar parameters
    should be specified as is,
    while lists and objects should
    use JSON format. Note that
    spaces for string values are
    not stripped. The output is
    not expected to be valid XML
    and is parsed with regular
    expressions.
46
47
48
49 Here are the functions available
    in JSONSchema format:
50
51 ## Server name: browsing_agent
52 ### Tool name: search_and_browse
53 Description: This tool is an agent
    that performs the subtask of
    searching and browsing the web
    for specific missing
    information and generating the
    desired answer. The subtask
    should be clearly defined,
    include relevant background,
    and focus on factual gaps. It
    does not perform vague or
    speculative subtasks.
54 Args:
55   subtask: the subtask to be

```

```

    performed.
56 Returns:
57   the result of the subtask.
58 Input JSON schema: {'properties':
    {'subtask': {'title': 'Subtask
    ', 'type': 'string'}}, '
    required': ['subtask'], 'title
    ': 'search_and_browseArguments
    ', 'type': 'object'}
59
60
61
62 # General Objective
63
64 You accomplish a given task
    iteratively, breaking it down
    into clear steps and working
    through them methodically.
65
66 ## Task Strategy
67
68 1. Analyze the user's request and
    set clear, achievable sub-
    goals. Prioritize these sub-
    goals in a logical order.
69 2. Start with a concise, numbered,
    step-by-step plan outlining
    how you will solve the task
    before taking any action.
70 3. Work through these sub-goals
    sequentially. After each step,
    adjust your plan as needed.
71 4. Use tools strategically to
    accomplish each sub-goal.
72 5. Revise earlier steps if new
    information emerges.
73
74 ## Tool-Use Guidelines
75
76 1. Each step must involve a single
    tool call, unless the task is
    already solved.
77 2. Before each tool call:
78   - Summarize what is known.
79   - Identify what is missing.
80   - Choose the most relevant tool
    .
81   - Verify all required
    parameters.
82 3. All tool queries must include
    full context.
83 4. Avoid vague queries. Each call
    should retrieve actionable
    information.
84 5. Extract and summarize partial
    information if a tool result
    is incomplete.
85
86 ## Tool-Use Communication Rules
87
88 1. Do not include tool results in
    your response.
89 2. Do not present the final answer

```

until the entire task is complete.

90 3. Do not mention tool names.

91 4. Do not engage in unnecessary back-and-forth.

92 5. Do not use non-existent tools.

93 6. Respond in the same language as the user's message.

94 7. If the task does not require tool use, answer directly.

95

96

97 # Agent Specific Objective

98

99 You are a task-solving agent that uses tools step-by-step to answer the user's question. Your goal is to provide complete, accurate and well-reasoned answers using additional tools.

100

101

102 user

103

104 On Cornell Law School website's legal information institute, under the fifth section of federal rules alphabetically, what word was deleted in the last amendment to the first rule in the article that has "witnesses" in the most titles as of 2021? You should follow the format instruction in the request strictly and wrap the final answer in \boxed{}

105

106

107 assistant

108

109 <think>

110 Okay, let's try to figure out this question step by step. The user is asking about a specific change on the Cornell Law School's Legal Information Institute website. They want to know which word was deleted in the last amendment to the first rule in an article that has "witnesses" in the most titles. This article is under the fifth section of the federal rules alphabetically, and they're looking at the status as of 2021.

111

112 First, I need to break down the question into manageable parts. Let me start by identifying the fifth section of the

federal rules alphabetically. The federal rules are organized alphabetically, so the fifth section would be the fifth letter in the alphabet. Let's see: A, B, C, D, E. So the fifth section is "E" rules. The Federal Rules of Evidence are under the E section because they start with "E" (Federal Rules of Evidence). So that's probably the section they're referring to.

113

114 Next, the article that has "witnesses" in the most titles. Looking at the Federal Rules of Evidence, the main article would be the Federal Rules of Evidence itself. The titles of the rules might include "witnesses" in their headings. For example, Rule 601 is about competency of witnesses, Rule 602 about lay witnesses, etc. So the Federal Rules of Evidence article likely has multiple rules with "witnesses" in their titles.

115

116 Now, the first rule in this article. The first rule in the Federal Rules of Evidence is Rule 601, which deals with the competency of witnesses. The question is about the last amendment to this rule as of 2021. I need to check when Rule 601 was last amended before 2021. Wait, but the user said "as of 2021," so I should consider the amendments up to that year.

117

118 I need to find out what changes were made to Rule 601 in its last amendment before 2021. The original Rule 601 was amended in 2011, I believe. Let me check: the 2011 amendments to the Federal Rules of Evidence included changes to several rules, including Rule 601. The specific change was the deletion of the word "competent" from the rule. Originally, the rule stated that a witness must be competent, but the amendment removed that term, effectively eliminating the competency

119 requirement except for certain
120 exceptions.

120 Wait, but I should confirm this.
If the last amendment to Rule
601 before 2021 was in 2011,
and the change was the
deletion of the word "
competent," then that would be
the answer. However, I need
to make sure that this is
accurate and that there haven'
t been any further amendments
after 2011 up to 2021. I
should verify if there were
any other amendments to Rule
601 after 2011 but before
2021. For example, the 2018
amendment to the Federal Rules
of Evidence included some
changes, but I need to check
if Rule 601 was affected.

121
122 Alternatively, maybe the user is
referring to a different rule.
Wait, the first rule in the
article is Rule 601. Let me
check the exact text of the
amendments. The 2011 amendment
to Rule 601 deleted the word
"competent" from the
definition, making the rule
about the qualifications of
witnesses, not their
competency. This was a
significant change. So the
word "competent" was deleted.

123
124 But I need to confirm this through
the Cornell Law website.
Since I can't access the
internet, I'll have to rely on
my existing knowledge.
However, to be thorough, I
should use the tool provided
to search the Cornell Law
website.

125
126 So the steps would be:

- 127
- 128 1. Identify the fifth section of
the federal rules
alphabetically, which is "E" (
Federal Rules of Evidence).
- 129 2. Find the article under this
section that has the most
titles with "witnesses" in
them. That would be the
Federal Rules of Evidence
themselves, particularly Rules
601-615.
- 130 3. Locate the first rule in this
article, which is Rule 601.

- 131 4. Check the last amendment to
Rule 601 before 2021.

- 132 5. Determine which word was
deleted in that amendment.

133
134 Using the tools available, I
should start by searching the
Cornell Law website for the
Federal Rules of Evidence,
specifically Rule 601. Then
check the amendments to this
rule as of 2021. The tool can
help get the exact text and
amendments.

135 </think>

136

137 <use_mcp_tool>

138 <server_name>browsing_agent</
server_name>

139 <tool_name>search_and_browse</
tool_name>

140 <arguments>

141 {

142 "subtask": "Identify the fifth
section of the federal rules
alphabetically on the Cornell
Law School's Legal
Information Institute website
."

143 }

144 </arguments>

145 </use_mcp_tool>

146

147

148 user

149

150 The fifth section of the federal
rules alphabetically on the
Cornell Law School's Legal
Information Institute website
is **"Federal Rules of
Bankruptcy Procedure."**

151

152 **### Supporting Information:**

153 The webpage at [<https://www.law.cornell.edu/rules>] (<https://www.law.cornell.edu/rules>) lists
the federal rules in the
following alphabetical order:

154 1. Federal Rules of Appellate
Procedure

155 2. Federal Rules of Civil
Procedure

156 3. Federal Rules of Criminal
Procedure

157 4. Federal Rules of Evidence

158 5. ****Federal Rules of Bankruptcy
Procedure****

159 6. U.S. Supreme Court Rules

160

161 This ordering is explicitly
presented on the site,
confirming the fifth entry as

the Bankruptcy Procedure rules
 162 .
 163
 164 assistant
 165
 166 <think>
 167 Okay, let's see. The user asked
 about the fifth section of the
 federal rules alphabetically
 on the Cornell Law website.
 The previous response
 correctly identified it as the
 Federal Rules of Bankruptcy
 Procedure. But now, the next
 step is to find the article
 under this section that has "
 witnesses" in the most titles.
 168
 169 Wait, the Federal Rules of
 Bankruptcy Procedure might not
 have many rules with "
 witnesses" in their titles.
 The user mentioned "the
 article that has 'witnesses'
 in the most titles as of
 2021." Maybe I need to check
 if there's a specific rule or
 article under Bankruptcy
 Procedure that references
 witnesses. Alternatively,
 perhaps there's a mix-up here.
 The initial assumption that
 the fifth section is
 Bankruptcy Procedure is
 correct based on the given
 data, but maybe the user
 intended a different section,
 like Evidence, which is fourth
 . However, the tool response
 confirmed it's fifth as
 Bankruptcy.
 170
 171 Alternatively, maybe the user made
 a mistake in the order, but I
 have to proceed with the
 information given. So, within
 the Federal Rules of
 Bankruptcy Procedure, I need
 to find rules with "witnesses"
 in their titles. Let me think
 : Bankruptcy rules typically
 deal with bankruptcy
 proceedings, creditors,
 debtors, etc. Rules about
 witnesses might be less common
 here. However, perhaps there'
 s a rule related to testimony
 or evidence in bankruptcy
 cases. For example, Rule 3003
 or something similar? Not sure
 . Alternatively, maybe the
 user confused the section, but

I need to follow the tool's
 response.
 172
 173 Alternatively, maybe the fifth
 section is actually the
 Federal Rules of Evidence, but
 the tool response said
 Bankruptcy. Wait, the tool
 response listed the order as:
 Appellate, Civil, Criminal,
 Evidence, then Bankruptcy. So
 the fifth is indeed Bankruptcy
 . But then, the article with "
 witnesses" would probably not
 be in Bankruptcy. This might
 be an error in the initial
 step. Wait, perhaps the user
 intended the fifth section as
 the Evidence rules, but the
 tool response said Bankruptcy.
 Let me double-check the
 alphabetization.
 174
 175 Federal Rules of Appellate
 Procedure (A), Civil (C),
 Criminal (C), Evidence (E),
 Bankruptcy (B). Wait,
 alphabetically, B comes before
 C. So the correct order
 should be: Appellate (A),
 Bankruptcy (B), Civil (C),
 Criminal (C), Evidence (E).
 But the tool response listed
 them as A, C, C, E, B. That
 seems incorrect. Maybe the
 tool's information is wrong,
 or the ordering on the site is
 different.
 176
 177 Alternatively, perhaps the site
 lists them in the order of the
 judiciary branches, not
 strictly alphabetical. However
 , the user said "
 alphabetically," so the fifth
 should be the fifth letter.
 This might be a mistake in the
 tool's response. But since
 the tool response is given, I
 need to proceed with that.
 However, this inconsistency
 could lead to error. If the
 fifth section is Bankruptcy,
 but the correct alphabetical
 order would have Bankruptcy as
 the second section, then the
 user might have intended a
 different section. This could
 be a critical point.
 178
 179 Alternatively, maybe the user
 meant the fifth section when
 considering all federal rules,

including those from different branches. But the initial search result from the tool seems to have ordered them as A, C, C, E, B. This is confusing. However, I must proceed with the given information. Therefore, under the fifth section (Bankruptcy), I need to look for an article with "witnesses" in the titles. If there's none, perhaps the initial step was incorrect, but the tool's response is taken as correct.

180

181 Alternatively, maybe the user's question has a mistake, and the intended section is Evidence (fourth), which would have rules about witnesses. But the tool response says fifth is Bankruptcy. This is a conflict. Given this, perhaps the correct approach is to check the Federal Rules of Evidence (fourth section) for rules with "witnesses" in titles and see if the user confused the section. However, strictly following the tool's response, the fifth is Bankruptcy.

182

183 Given this confusion, the next step is to check under the Bankruptcy rules for any mention of witnesses. If there are none, then perhaps the initial step was wrong, and the user intended a different section. But according to the tool's response, the fifth is Bankruptcy. So, I'll proceed under that and look for any rules in Bankruptcy Procedure with "witnesses" in their titles. If none, the answer might be that there's no such rule, but the user's question implies there is one. This suggests a possible error in the tool's response or the initial step.

184

185 Alternatively, perhaps the fifth section is actually the Federal Rules of Evidence, and the tool made a mistake in ordering. In that case, the correct fifth would be Evidence. But I need to follow the given tool response. This

is a dilemma. However, since the tool response is part of the process, I must proceed with the information provided, even if it might be incorrect. Therefore, I'll proceed to search within the Federal Rules of Bankruptcy Procedure for any rules with "witnesses" in their titles, even though it's unlikely. If no such rules are found, the answer might be that the premise is incorrect, but the user's question requires an answer based on the given steps.

186 </think>

187

188 <use_mcp_tool>

189 <server_name>browsing_agent</server_name>

190 <tool_name>search_and_browse</tool_name>

191 <arguments>

192 {

193 "subtask": "Find the article under the Federal Rules of Bankruptcy Procedure that has the most titles containing the word 'witnesses' as of 2021."

194 }

195 </arguments>

196 </use_mcp_tool>

197

198

199 user

200

201 The article under the Federal Rules of Bankruptcy Procedure that contains the term "witnesses" is **Rule 9014**, titled **"Contested Matters"**. This rule governs procedures for contested matters in bankruptcy proceedings and explicitly addresses the **testimony of witnesses** in disputed factual issues. Here's the detailed supporting information:

202

203 1. **Rule 9014 Overview**:

204 Rule 9014 outlines procedures for contested matters not otherwise governed by other rules. It mandates that relief be requested by motion, with notice and an opportunity for a hearing. It also specifies that testimony of witnesses

regarding disputed factual issues must be taken in the same manner as in an adversary proceeding (§ 9014(d)).

205

206 2. ****Direct Mention of Witnesses**

207 ****:**

208 The rule explicitly states:

208 > ***"Testimony of witnesses with**

respect to disputed material factual issues shall be taken in the same manner as testimony in an adversary proceeding."*****

209 This directly ties the procedural framework for witness testimony to adversarial proceedings under the Bankruptcy Code.

210

211 3. ****Contextual Relevance**:**

212 While no other rule in the 2021 edition of the Federal Rules of Bankruptcy Procedure explicitly includes the word "witnesses" in its title, Rule 9014 is the only rule that addresses witnesses in its content. It aligns with broader bankruptcy procedures where witness testimony may be required in contested matters (e.g., objections to claims, motions for relief).

213

214 4. ****Source Verification**:**

215 The content was retrieved from the official 2021 Federal Rules of Bankruptcy Procedure website:

216 <https://www.federalrulesofbankruptcyprocedure.org/part-ix/rule-9014/>

217

218 **### Conclusion:**

219 ****Rule 9014**** is the sole article under the Federal Rules of Bankruptcy Procedure (2021) that contains the term "witnesses" in its text, making it the answer to the query.

220

221 assistant

222

223 <think>

224 Okay, the user is asking about the

deletion of a word in the last amendment to the first rule of the article that has "witnesses" in the most titles. From the previous steps, we identified that the fifth section is the Federal Rules of Bankruptcy Procedure, and the relevant article is Rule 9014, which mentions "witnesses" in its text. However, Rule 9014 is not the first rule in the Bankruptcy Procedure section. Wait, the user mentioned the first rule in the article. So maybe I made a mistake here.

226

227 Wait, the user's original question refers to the first rule in the article that has "witnesses" in the most titles. In the previous steps, the assistant found that Rule 9014 is the only rule in the Bankruptcy Procedure section with "witnesses" in its text. But Rule 9014 isn't the first rule; it's part of Part IX. The first rule in the Bankruptcy Procedure is Rule 1001. However, Rule 1001 doesn't mention witnesses. So there's a contradiction here.

228

229 Alternatively, perhaps the user confused the sections, and the intended section was the Federal Rules of Evidence (fourth section), which does have multiple rules with "witnesses" in their titles. For example, Rules 601, 602, 603, etc., all relate to witnesses. If the fifth section was mistakenly identified as Bankruptcy, but the correct fifth section alphabetically is actually the Federal Rules of Evidence (since E comes after B), then the previous steps were incorrect.

230

231 This suggests a possible error in the initial step where the fifth section was identified as Bankruptcy. Alphabetically, the order should be: Appellate (A), Bankruptcy (B), Civil (C), Criminal (C), Evidence (E). Therefore, the fifth section is actually the

Federal Rules of Evidence, not Bankruptcy. The tool's response might have incorrectly listed the order, leading to confusion.

232

233 Given this, the user's question likely intended the Federal Rules of Evidence (fourth section) as the correct section. However, since the tool response said fifth is Bankruptcy, the assistant proceeded with that. Now, the current step is to check the first rule in Rule 9014 (Bankruptcy) for amendments. But Rule 9014 is part of the Bankruptcy Procedure, and the first rule in that section is Rule 1001. However, Rule 1001 doesn't mention witnesses. Therefore, there's inconsistency here.

234

235 Alternatively, perhaps the assistant should re-evaluate the initial step. If the fifth section is actually the Federal Rules of Evidence (E), then the article with "witnesses" in titles would be Rules 601-615. The first rule there is Rule 601, which deals with the competency of witnesses. The last amendment to Rule 601 before 2021 was in 2011, where the word "competent" was deleted. That would make sense, as the user's question seems to reference that change.

236

237 So, the mistake might be in the initial identification of the fifth section. If the fifth section is actually the Federal Rules of Evidence, then the correct path would be : Federal Rules of Evidence (E) as the fifth section, leading to Rule 601 as the first rule in the article with most "witnesses" in titles. The last amendment to Rule 601 deleted the word "competent" in 2011. Hence, the answer would be "competent".

238 </think>

239

240 \boxed{competent}